



Python for Fixed Income

COURSE 1 · PYTHON FOUNDATIONS

pandas: Series and DataFrames

WEEK 2 · DAY 2

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

By the end of today you can

Load a holdings file

Its size, its columns, their types, what is missing.

```
holdings = pd.read_csv(data_path)
```

Look up a bond

Columns by name, and a bond by its CUSIP.

```
by_cusip.loc['99039NAC0', 'price']
```

Filter, count, and rank

By sector, rating, price, spread, or duration.


```
holdings[holdings['sector'] == 'Energy']
```

Add a column, run a what-if

On a copy, without touching the desk's record.

```
what_if = by_cusip.copy()
```


From four bonds to two hundred



Each bond now comes with text (white) and numbers (dark)

CUSIP	Issuer name	Sector
Rating	Maturity date	Price
Spread	Duration	Position size

The tools so far, and the one for today



A Series is one labelled column

prices

Built from two lists. Looked up by label, or by position.

index	values
QVNE	98.50
DNMR	103.25
PLWK	95.00
HLVF	101.00

dtype: float64

```
prices = pd.Series(price_list, index=ticker_list)

print('PLWK:', prices['PLWK'])
PLWK: 95.0
print('First price:', prices.iloc[0])
First price: 98.5
```

Two arrays in different orders

Multiplied position by position. No error, and four wrong numbers.

Two Series in different orders

Lined up by ticker first. In Excel, the VLOOKUP on the ticker, done without being asked.

A DataFrame is a table of Series

desk: four rows, six columns, shape (4, 6)

	Ticker	Coupon	Maturity	Price	Face	Rating
0	QVNE	5.250	2033	98.50	1500000	BBB
1	DNMR	6.125	2034	103.25	500000	BB+
2	PLWK	4.750	2031	95.00	750000	A-
3	HLVF	7.500	2032	101.00	250000	B+
dtype	str	float64	int64	float64	int64	str

Each column has one type, like an array. The table as a whole can mix them.

Every column is a Series

```
desk_prices = desk['Price']  
print(type(desk_prices))  
<class 'pandas.Series'>
```

In Excel

A worksheet with a header row.

The column names are the headers.

The index is the row numbers down the left edge.

The holdings file: 200 bonds, 19 columns

```
holdings = pd.read_csv(data_path)
```

One row per bond. The first line of the file becomes the column names. Eight of the nineteen columns are shown.

	cusip	ticker	sector	rating	coupon	price	oas	par_held
0	99080CAE8	BEZE	Financials	BBB-	7.000	106.727	147	5000000
1	99049XAE2	BEZH	Industrials	B	8.125	97.556	535	250000
2	99039NAC0	BEZL	Transportation	BB	5.125	98.730	153	1500000
units	text	text	text	text	percent of face	per 100 of face	basis points	dollars of face

holdings.shape

(200, 19): 200 bonds, 19 columns.

holdings.info()

200 non-null in every column. Nothing is missing.

```
holdings['oas']  
# one name: a Series  
holdings[['ticker', 'coupon', 'maturity', 'price']]  
# a list of names: a smaller table
```

Rows: `iloc` by position, `loc` by label

```
by_cusip = holdings.set_index('cusip')
```

position	cusip (the index)	ticker	coupon	price
0	99080CAE8	BEZE	7.000	106.727
1	99049XAE2	BEZH	8.125	97.556
2	99039NAC0	BEZL	5.125	98.730

Sort the table or remove a bond, and position 0 is a different bond. A label stays with its row.

`iloc` By position. A slice leaves out its end.

```
by_cusip.iloc[0:3]  positions 0, 1, 2
```


`loc` By label. A slice includes its end.

```
by_cusip.loc['99080CAE8':'99039NAC0']  both ends
```

Question	pandas	Excel
Price of a CUSIP	<code>by_cusip.loc['99039NAC0', 'price']</code>	<code>=VLOOKUP("99039NAC0", B:I, 8, FALSE)</code>
A CUSIP that is not there	<code>KeyError</code>	<code>#N/A</code>

A column is a Series, so it can be summarized

`holdings['sector'].value_counts()`



A count of an invented portfolio. In Excel, a pivot table.

The same job as the Excel function, with the column named

pandas	Excel
<code>holdings['market_value'].sum()</code>	<code>=SUM(S2:S201)</code>
<code>holdings['oas'].mean()</code>	<code>=AVERAGE(M2:M201)</code>
<code>holdings['oas'].median()</code>	<code>=MEDIAN(M2:M201)</code>
<code>holdings['price'].min()</code>	<code>=MIN(I2:I201)</code>

115 `nunique()`

tickers. Many issuers have more than one bond.

200 `nunique()`

CUSIPs for 200 rows. Every label is different.

A filter is a column of True and False

holdings

	ticker	sector
0	BEZE	Financials
1	BEZH	Industrials
...	...	...
198	ZEPN	Energy
199	ZEPN	Energy

the mask

is_energy
False
False
...
True
True

holdings[is_energy]

	ticker	sector
67	GALN	Energy
68	GALN	Energy
...	...	...
198	ZEPN	Energy
199	ZEPN	Energy

17

Energy bonds: the mask, summed

38,830,005

Their total market value

Question	pandas	Excel
Show the Energy bonds	<code>holdings[holdings['sector'] == 'Energy']</code>	AutoFilter on the sector column
Count the Energy bonds	<code>(holdings['sector'] == 'Energy').sum()</code>	<code>=COUNTIF(E:E, "Energy")</code>
Market value of the Energy bonds	<code>holdings.loc[is_energy, 'market_value'].sum()</code>	<code>=SUMIF(E:E, "Energy", S:S)</code>

& narrows. | widens.

& keeps the rows where both conditions are True

```
long_and_low = (holdings['duration'] > 10) & (holdings['price'] < 90)
print('Count:', long_and_low.sum())
Count: 6
```

6 bonds pass both tests

| keeps the rows where at least one condition is True

```
low_or_wide = (holdings['price'] < 80) | (holdings['oas'] > 1000)
print('Count:', low_or_wide.sum())
Count: 4
```

4 bonds pass at least one test

`( ... ) & ( ... )`

Each condition needs its own round brackets.

and, or

ValueError: The truth value of a Series is ambiguous.

`isin(['A+', 'A', 'A-'])`

True where the value is in the list. 56 single-A bonds.

Sorting hands back a sorted copy

```
widest = holdings.sort_values('oas', ascending=False)
```

	ticker	sector	rating	price	oas
65	FESR	Technology	CCC+	87.116	1551
17	CLDB	Consumer	B	62.045	1446
132	PRAN	Financials	CCC+	89.407	871
110	NEVN	Utilities	B-	96.616	720
111	NEVN	Utilities	B-	94.053	719

The labels moved with their rows: the first row is still labelled 65. holdings itself is not sorted. The issuers are invented.

Sort by rating, and text sorts letter by letter

A	A+	A-	AA	AA+	AA-	AAA
1	2	3	4	5	6	7

Alphabetical order, not credit order. AAA comes out seventh.

```
holdings.nlargest(5, 'market_value')
```

A shortcut for sort, then head.

In Excel: Data, Sort

Excel sorts the sheet in place. pandas leaves the original order alone.

A new column is a formula column

```
holdings['coupon_income'] = holdings['par_held'] * holdings['coupon'] / 100
```

	ticker	coupon	par_held	coupon_income
0	BEZE	7.000	5000000	350000.0
1	BEZH	8.125	250000	20312.5
2	BEZL	5.125	1500000	76875.0

In Excel

=R2*G2/100

A formula in a new column, filled down. A pandas column holds values, not live formulas: change an input, and rerun the line.

1 Work on a copy

```
what_if = by_cusip.copy()
```

2 Set one value with loc


```
what_if.loc['99039NAC0',  
            'price'] = 97.0
```

3 Recalculate, then compare

```
print('Change in market value:',  
      round(change, 2))  
Change in market value: -25950.0
```

The AI check

Every notebook ends with code written in the style of an AI assistant. It has one bug. Find it.



Code that runs is not the same as code that is right.

What you can do now

- Load a holdings file and take a first look at its size, columns, and types
- Pull out the columns a question needs, and look a bond up by its CUSIP
- Filter by sector, rating, price, spread, or duration, and count what passes
- Rank the holdings: widest spreads, largest positions
- Add a calculated column, and run a what-if on a copy

Tomorrow pandas: joining tables, reading and saving files, summaries by group, and dates