



Python for Fixed Income

COURSE 1 · PYTHON FOUNDATIONS

Case Study: Consumer Credit

WEEK 4 · DAY 1

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

The problem and the data

Real public data: the UCI Default of Credit Card Clients dataset

30,000

credit card accounts

25

columns in the file

6

months of records

The first three accounts, seven of the 25 columns. Records from April to September 2005.

ID	LIMIT_BAL	SEX	EDUCATION	AGE	PAY_0	default payment next month
1	20000	2	2	24	2	1
2	120000	2	2	26	-1	1
3	90000	2	2	34	0	0

```
raw = pd.read_csv(data_folder + 'credit_card_default.csv')
print(raw.shape)  # (30000, 25)
```

Amounts are in New Taiwan dollars (NT dollars). PAY_0 is September 2005, back to PAY_6, April. There is no PAY_1.

Source: Yeh, I. (2009). Default of Credit Card Clients [Dataset]. UCI Machine Learning Repository. <https://doi.org/10.24432/C55S3H>.

License: Creative Commons Attribution 4.0 International (CC BY 4.0), <https://creativecommons.org/licenses/by/4.0/legalcode>

Changes made for this course: The original is an Excel file with two header rows. The first, which holds the labels X1 to X23 and Y, was dropped, and the sheet was saved as credit_card_default.csv. Nothing else was changed: all 30,000 rows and 25 columns are there with their original names and values.

The colleague asks

What does this file show about which accounts defaulted the following month?

The answer is a description: rates in this file, each with its count. It does not say why an account defaulted, and it is not advice on lending.

Taiwan, 2005

The file is one set of credit card accounts, in Taiwan, in 2005. Consumer credit markets differ from one country and one period to another, including from US credit. Nothing found here carries over to another market.

The file against its documentation

Column	The documentation lists	The file also holds	Accounts holding them	Decision
SEX	1, 2	nothing else	0	Nothing to decide
EDUCATION	1, 2, 3, 4	0, 5, 6	345	Kept, labelled not documented
MARRIAGE	1, 2, 3	0	54	Kept, labelled not documented
PAY_0, September status	-1, and 1 to 9	-2, 0	17,496	Kept as code -2 and code 0
default payment next month	1, 0	nothing else	0	Renamed default in the working copy

A code the documentation does not define is kept, counted, and called not documented. It is not dropped or guessed at.

```
education_labels = {1: 'graduate school', 2: 'university',
                    3: 'high school', 4: 'others'}
cards['education'] = cards['EDUCATION'].map(education_labels)
print(cards['education'].isna().sum()) # 345
cards['education'] = cards['education'].fillna('not documented')
```

A name with a space cannot follow the dot

`raw.default payment next month.mean()` stops with `SyntaxError: invalid syntax`.
`raw['default payment next month']` works.

`map()` left 345 blanks and raised no error. `isna().sum()` caught it.

Nothing is missing. 35 rows repeat once the ID is left out: kept.

The same in Excel: sheet cards, rows 2 to 30001, EDUCATION in column D

`=COUNTIF(D2:D30001, 5)` returns 280 `=COUNTIF(D2:D30001, "<1") + COUNTIF(D2:D30001, ">4")` returns 345
`=IFERROR(VLOOKUP(D2, $AB$2:$AC$5, 2, FALSE), "not documented")` the label. Codes in AB2 to AB5, meanings in AC2 to AC5.

Every code at once: a pivot table with EDUCATION in Rows and Count of ID in Values. It has seven rows.

The outcome: a rate and its count

The outcome holds 1 for a default and 0 for none

6,636

defaults: the sum

/

30,000

rows: the count

=

0.2212

the mean: a rate

```
print(cards['default'].mean())    # 0.2212
defaults = cards['default'].sum() # 6636
accounts = len(cards)            # 30000
```

The mean of a 0/1 column is the share of rows that hold 1.

22.12 percent of the accounts in this file defaulted the following month.

The same in Excel, with the outcome in column Y

=SUM(Y2:Y30001)	returns 6,636, the defaults
=COUNT(Y2:Y30001)	returns 30,000, the accounts
=AVERAGE(Y2:Y30001)	returns 0.2212. As a percentage: 22.12%

A rate needs its count

Group	Defaults	Accounts	Rate, percent
EDUCATION code 0	0	14	0.00
All accounts	6,636	30,000	22.12


0 defaults among 14 accounts. 1 default in 14 would be 7.14 percent. On 30,000 accounts, one account more or less moves the rate by 0.003 of a percentage point.

```
code_zero = cards[cards['EDUCATION'] == 0]
print(code_zero['default'].mean() * 100) # 0.0
print(len(code_zero))                   # 14
```

From here on, no rate is shown without the number of accounts behind it.

One column at a time

Credit limit, one shape for each outcome



51.4 percent of the 6,636 accounts that defaulted had a limit below 100,000 NT dollars, and 34.4 percent of the 23,364 that did not. Shares of a group, not default rates.

In Excel: limit in B, bill in M, payment in S

=AVERAGE(B2:B30001) returns 167,484

=MEDIAN(B2:B30001) returns 140,000

=SKEW(S2:S30001) returns 14.67

Chart: Insert, Insert Statistic Chart, Histogram.

Four columns: amounts in NT dollars, age in years

Column	What it is	Median	Mean	Skew	Lowest	Highest
LIMIT_BAL	credit limit	140,000	167,484	0.99	10,000	1,000,000
AGE	age	34	35.5	0.73	21	79
BILL_AMT1	September bill	22,382	51,223	2.66	-165,580	964,511
PAY_AMT1	paid in September	2,100	5,664	14.67	0	873,552

81 different limits. Most common: 50,000, on 3,365 accounts. 206 are above 500,000.

The September bill: what the documentation does not explain

590 bills are below zero, 2,008 are exactly zero, 2,115 are above the limit.

Kept, counted, and reported as they are. No guess is made.

5,249 accounts paid nothing in September, and 471 paid more than 50,000. With a skew of 14.67, quote the median and the quartiles.

```
sns.histplot(data=cards, x='limit_thousands', hue='outcome',
              stat='percent', common_norm=False, element='step')
print((cards['BILL_AMT1'] > cards['LIMIT_BAL']).sum()) # 2115
```


=COUNTIF(M2:M30001, "<0") returns 590

=COUNTIF(S2:S30001, 0) returns 5,249

=SUMPRODUCT(--(M2:M30001 > B2:B30001)) returns 2,115, the bills above the limit

The outcome against one column

Default rate by education label, with a reference line



Counts are not rates: the group with more defaults had the lower rate.

A bar on 123 accounts looks as solid as one on 14,030: read the counts.

```
by_sex = cards.groupby('sex')['default'].agg(['count', 'sum', 'mean'])
by_sex.columns = ['accounts', 'defaults', 'rate_pct']
```

In Excel: SEX in column C, the outcome in column Y

<code>=COUNTIF(C2:C30001, 1)</code>	returns 11,888, the accounts
<code>=SUMIF(C2:C30001, 1, Y2:Y30001)</code>	returns 2,873, the defaults
<code>=AVERAGEIF(C2:C30001, 1, Y2:Y30001)</code>	returns 0.2417, the rate


Column	Label	accounts	defaults	rate_pct
sex	female	18,112	3,763	20.78
sex	male	11,888	2,873	24.17
education	graduate school	10,585	2,036	19.23
education	university	14,030	3,330	23.73
education	high school	4,917	1,237	25.16
education	others	123	7	5.69
education	not documented	345	26	7.54
marriage	married	13,659	3,206	23.47
marriage	single	15,964	3,341	20.93
marriage	others	323	84	26.01
marriage	not documented	54	5	9.26

Differences between groups in this file do not explain themselves, and they are not a basis for treating people differently.

In each grouping the accounts add up to 30,000.

Bands: `pd.cut()` and `pd.qcut()`

Default rate by credit limit band, with the reference line



Each band has more than 7,000 accounts: no rate rests on a thin group.

An association. The documentation does not say how the limits were set.

`pd.qcut(cards['PAY_AMT1'], 10)`: `ValueError`, Bin edges must be unique. 5,249 accounts paid nothing, so two edges are both 0.

`=PERCENTILE.INC(B2:B30001, 0.25)` returns 50,000

`=AVERAGEIFS(Y2:Y30001, B2:B30001, ">50000", B2:B30001, "<=140000")` returns 0.2472

`=COUNTIFS(B2:B30001, ">50000", B2:B30001, "<=140000")` returns 7,614

`=PERCENTILE.INC(S2:S30001, 0.1)` returns 0, the same as the minimum payment

`=VLOOKUP(B2, $AH$2:$AI$5, 2, TRUE)` returns the band

Lowest limit of each band in AH2 to AH5, band names in AI2 to AI5.

limit_band	accounts	defaults	rate_pct
10,000 to 50,000	7,676	2,440	31.79
50,001 to 140,000	7,614	1,882	24.72
140,001 to 240,000	7,643	1,326	17.35
240,001 to 1,000,000	7,067	988	13.98

`pd.qcut()`: the data chooses the edges, at the quartiles

```
limit_band = pd.qcut(cards['LIMIT_BAL'], 4)
```


7,676, 7,614, 7,643, and 7,067 accounts, not 7,500 each: the 3,365 accounts at exactly 50,000 stay in one band.

`pd.cut()`: you choose the edges

```
pd.cut(cards['AGE'], bins=[20, 30, 40, 50, 60, 80], labels=age_names)
```

Five age bands: 11,013, 10,713, 6,005, 1,997, and 272 accounts.

Repayment status in September



PAY_0	The documentation	accounts	defaults	rate_pct
-2	does not define it	2,759	365	13.23
-1	pay duly	5,686	954	16.78
0	does not define it	14,737	1,888	12.81
1	delay of one month	3,688	1,252	33.95
2	delay of two months	2,667	1,844	69.14
3	delay of three months	322	244	75.78
4	delay of four months	76	52	68.42
5	delay of five months	26	13	50.00
6	delay of six months	11	6	54.55
7	delay of seven months	9	7	77.78
8	delay of eight months	19	11	57.89

```
by_status = rate_table('PAY_0')
```

Codes -2, -1, and 0 are below the line. Every code from 1 up is above it.
Codes -2 and 0: the two lowest rates, and 17,496 accounts between them.



What they mean is not in the documentation, and the rates do not say.

```
=COUNTIF(G2:G30001, -2) + COUNTIF(G2:G30001, 0) returns 17,496: the accounts with code -2 or code 0 in September
```

Codes 4 to 8 hold 141 accounts between them: thin groups.

Read each rate with its count. The table is part of the chart.

Two columns: the rate and the count



```
rate_pivot = pd.pivot_table(cards, values='default', index='status_group',  
                             columns='limit_band', aggfunc='mean')
```

The highest rate, 83.9 percent, is on the smallest cell: 31 accounts.


Status 2, and 3 to 8: 67.5 to 83.9 percent in every limit band. Code 0: 19.2 to 7.4 percent. 24 cells, 31 to 4,288 accounts, 30,000 in all.

`=COUNTIFS(G2:G30001, 2, B2:B30001, "<=50000")` returns 1,049 `=AVERAGEIFS(Y2:Y30001, G2:G30001, 2, B2:B30001, "<=50000")` returns 0.684

Pivot table: status group in Rows, limit band in Columns, outcome in Values by Average or Count. Colours: Conditional Formatting, Color Scales.

Derived columns

Months with a payment delay: a status code of 1 or more



No month of delay: 11.71 percent of 19,931 accounts defaulted.

All six months: 70.32 percent of 1,341 accounts defaulted.

By outcome: 35.2 percent of the 6,636 accounts that defaulted had no month of delay. A percentage is a share of some total. Say which.

```
cards['months_delayed'] = (cards[pay_columns] >= 1).sum(axis=1)
```

=M2 / B2 0.196 for the first account

=COUNTIF(G2:L2, ">=1") 2 for the first account

=MEDIAN(AK2:AK30001) returns 0.061

=S2 / M2 #DIV/0! on the 2,008 rows where M is 0

=IF(M2 > 0, S2 / M2, "") blank on 2,598 rows. AK is that helper column.

Utilisation: September bill / credit limit. Median 0.314.

utilisation_band	accounts	defaults	rate_pct
zero or below	2,598	643	24.75
above 0 to 0.25	11,440	1,923	16.81
above 0.25 to 0.5	3,588	780	21.74
above 0.5 to 0.75	3,579	929	25.96
above 0.75 to 1	6,680	1,725	25.82
above 1	2,115	636	30.07


```
cards['utilisation'] = cards['BILL_AMT1'] / cards['LIMIT_BAL']
```

Share of the bill paid: look at the denominator

(cards['PAY_AMT1'] / cards['BILL_AMT1']).mean() prints inf, with no error: 2,008 bills are exactly zero. On the 27,402 bills above zero the median share is 0.061 and the mean 1.776. Quote the median.

Correlation: a first sort

The six monthly bills: 0.80 or more with each other



One month apart: 0.92 to 0.95. Six columns, nearly the same thing.

```
bill_corr = cards[bill_columns].corr()  
with_outcome = cards[columns_to_check].corr()['default']
```

The same in Excel

=CORREL(M2:M30001, N2:N30001) returns 0.95

=CORREL(B2:B30001, Y2:Y30001) returns -0.154

Colours: Home, Conditional Formatting, Color Scales.

Correlation with the outcome

Column	Correlation
months_delayed	0.398
PAY_0	0.325
utilisation	0.086
AGE	0.014
BILL_AMT1	-0.020
PAY_AMT1	-0.073
LIMIT_BAL	-0.154

Age: the rates behind 0.014

age_band	accounts	rate_pct
21 to 30	11,013	22.44
31 to 40	10,713	20.43
41 to 50	6,005	23.30
51 to 60	1,997	25.24
61 to 79	272	26.84

Down, then up: a straight line fits a dip badly, so the correlation is near zero.

What a correlation with a 0/1 outcome can and cannot say

Can: give the direction of a straight-line association, in one number.

Cannot: say what the default rate was in any group. The table does.

Cannot: see a pattern that is not a straight line, as with age.

Treats a code as an amount: PAY_0 as evenly spaced numbers.

Says nothing about cause.

A first sort. The rate tables, with their counts, are the findings.

The written observations


Each line is a rate in this file with the number of accounts behind it. None is a statement of cause or a basis for a lending decision.

1. The default rate	6,636 of the 30,000 accounts defaulted the following month: 22.12 percent.
2. September status	A delay of two months, code 2: 69.14 percent of 2,667 accounts. A delay of one month, code 1: 33.95 percent of 3,688. Pay duly, code -1: 16.78 percent of 5,686.
3. Months with a delay	No month of delay in the six: 11.71 percent of 19,931 accounts. A delay in all six: 70.32 percent of 1,341.
4. Credit limit	50,000 NT dollars or less: 31.79 percent of 7,676 accounts. Above 240,000: 13.98 percent of 7,067. Within a September status of 2: 67.5 to 71.1 percent in all four limit bands.
5. Utilisation	Bill above zero and up to a quarter of the limit: 16.81 percent of 11,440 accounts. Bill above the limit: 30.07 percent of 2,115.
6. Demographic columns	Sex: 24.17 percent (male, 11,888 accounts), 20.78 (female, 18,112). Education: 19.23 (graduate school, 10,585), 23.73 (university, 14,030), 25.16 (high school, 4,917). Marital status: 23.47 (married, 13,659), 20.93 (single, 15,964). Age band: 20.43 (31 to 40, 10,713), 26.84 (61 to 79, 272).
7. What is left open	17,496 accounts hold a September code, -2 or 0, that the documentation does not define: 13.23 percent (code -2, 2,759 accounts) and 12.81 percent (code 0, 14,737). Any reading of this file depends on what those codes mean.

`cards.groupby(column)['default'].agg(['count', 'sum', 'mean'])` is behind each line. In Excel: COUNTIF, SUMIF, and AVERAGEIF on the same condition.

The AI check

Every notebook ends with code written in the style of an AI assistant. It has one bug. Find it.



Code that runs is not the same as code that is right.

What you can do now

- Check a file against its documentation, and keep the codes it does not define
- Turn a 0/1 column into a rate, and show its count beside it
- Find the rate for each group, and for bands made with `pd.cut` and `pd.qcut`
- Read a table of rates beside its table of counts
- Write up what a file shows, and stop where the file stops

Week 4, Day 2 Analyzing Text. Regular expressions, cleaning, and turning text into numbers, on FOMC statements