



Python for Fixed Income

COURSE 4 · MACHINE LEARNING

PCA and t-SNE

WEEK 3 · DAY 4

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

From numbers you chose to numbers the data choose

Course 3

Shapes you chose: 2s10s, 5s30s, the 2s5s10s fly. Weights a desk picked.

Day 11

Tenors so correlated a regression could hardly tell them apart.

Today: PCA

A few weighted sums that carry most of how the curve moved. No target.

Real Treasury data, January 2015 to October 2026. A component is a shape of movement in that history; a score is how much of that shape a past date had.

No component says what rates will do, or what anyone should do. Nothing after the last row.

The title says what the method does


Karl Pearson, November 1901

"On lines and planes of closest fit to systems of points in space." The Philosophical Magazine, volume 2, issue 11, pages 559 to 572.

Today the points are 2,940 market dates, each a point in eleven dimensions, one per tenor.

Source, read 4 October 2026: the Crossref record of the paper (api.crossref.org), DOI 10.1080/14786440109462720. Bibliographic facts only; the paper itself was not read for this course.

Many columns that move together



0.98

20 Yr with 30 Yr daily changes


0.43

1 Yr with 30 Yr, the furthest pair

2,939 daily changes in bp of the 1 Yr to 30 Yr: eight columns, far fewer ways of moving.

Correlated columns carry less than their count.

PCA on levels: loadings and their shapes




	Share	Shape
PC1	95.62%	level
PC2	3.38%	slope
PC3	0.86%	curvature

Three components: 99.86 percent of the variance of eleven yields.

Level, slope, curvature: names for shapes of weights in this history.

PCA on daily changes



Score	Course 3 measure	Corr
PC1	average change	0.9995
PC2	change in 2s10s	-0.9174
PC3	change in 2s5s10s	-0.4509


Shares: 85.04, 11.02, 2.11 percent (98.17 in three).

PC3 is mostly the 1 Yr (loading 0.796). The loadings, not the name, say what it is.

```
signs = np.where(loadings.sum(axis=0) < 0, -1.0, 1.0)
```

The sign rule: loadings of each component sum to a positive number.

Scores over time, as history



Score	Lowest	Highest
Level	2020-08-04	2023-10-19
Slope	2023-05-04	2026-09-30

Allowed names

- level score
- slope score
- curvature score

Never called

- a forecast
- a view on rates
- a signal

Each one a score of a past date, inside the file's span.

PCA as a feature step, and lookahead

The no-lookahead check, t = 2021-01-04

Fitted up to t, later rows removed: unchanged, True.

Fitted on the whole history: the same scores? False.

Walk-forward, 19 folds, 2,394 test days: the 30 Yr change given the same day's other changes.

```
Pipeline([("pca", PCA(n_components=3)),  
          ("line", LinearRegression())])
```

0.9 bp

three components, then a line

0.7 bp

the line on all seven tenors

When PCA builds features it goes inside a pipeline, refitted on each fold's training rows.

Scaling before PCA, on mixed units


Column	Std, own units	PC1 raw	PC1 scaled
log_oas	0.735	0.179	0.602
duration	3.251	-0.252	-0.005
coupon	1.654	0.317	0.542
score	3.431	0.897	0.586

	Raw	Scaled
PC1 share	54.46	64.96
PC2 share	40.72	25.70

Unscaled, PC1 is mostly the column with the largest spread in its own units. Measure duration in months and it would take over.

Different units: scale first. The bonds are invented; this describes the course's generator.

t-SNE: a picture of many columns



Can be trusted for

Which rows sit near which; a first look at whether there are groups.

Cannot be trusted for

Distances between groups, group sizes, the axes, anything printed.

Excel side by side

Excel side by side (tested in Excel)

`=(B3-B2)*100` (K3, fill right and down)

`=CORREL(10 Yr changes, 20 Yr changes)`

`=COVARIANCE.S(2 Yr changes, 10 Yr changes)`

`=MMULT(K3:R2941-BYCOL(K3:R2941,LAMBDA(c,AVERAGE(c))),loadings!B2:D9)`

`=MMULT(K3:R2941,loadings!B2:D9)`

`=MMULT(covariance, PC1 loadings)/PC1 loadings`

`=F2/H2, with =SUMPRODUCT(covariance*MUNIT(8)) in H2`

the changes in bp, exactly

0.9529

20.9902 bp squared

the scores

not the scores: off by 0.333 bp

181.4263 on all 8 rows

PC1 share 85.04%

Excel has no PCA and no t-SNE. It confirms components it is given; it cannot find them. Centre first.

Save the day: a tag

1. Module

pca_loadings and explained_table, three tests

2. Results

results.csv: the walk-forward pipeline, 4 decimals

3. Tag

git tag week3-teaching


```
git log --oneline --decorate
```

HASH (HEAD -> main, tag: week3-teaching) Experiment log: walk-forward test metrics of three curve components then a line

A tag names one commit and never moves: the module as it stood at the end of the teaching days.

The AI check

The docstring is the prompt. The code that comes back runs, and has one modelling bug.



Today: PCA fitted on the whole history, then used for the first fold: 0.8 bp against 0.8 bp. Only the test catches it.

Code that runs is not the same as code that is right.

What you can do now

- Three components: 99.86% of levels, 98.17% of daily changes
- Level, slope, curvature: names for shapes, where the loadings show them
- Scores are the centred data times the loadings; the sign rule fixes the sign
- PCA that builds features goes inside a pipeline; the test proves it
- Scale first on mixed units; t-SNE is charted, never measured

Week 3, Day 5 The capstone: three models on a real file of company ratios