



Python for Fixed Income

COURSE 5 · ADVANCED MACHINE LEARNING

LightGBM and Stacking

WEEK 2 · DAY 4

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

Three boosters, then a stack

Three boosters

LightGBM

leaf-wise trees, bins

HistGradientBoosting

scikit-learn's own, bins

XGBoost

day 8, depth-limited trees

A stack

Base models give a probability for every account. A small model on top takes them as columns and learns how much to trust each.

The trap: which outputs the top model learns from.

Can several models that score about the same be combined into one that beats them all?

LightGBM: leaves, not depth, and bins

Depth-limited (day 8)

Splits every node at one level before going deeper. `max_depth=4` allows at most 16 leaves.

Leaf-wise (LightGBM)

Splits whichever leaf gains most, anywhere, until `num_leaves` leaves (31 by default).

31

leaves in each of 300 trees

8 to 22

questions to the deepest leaf

14,404 to 207

BILL_AMT1 values to split points

At most 255 bins a column (`max_bin`): never more than 254 split points to weigh.

LightGBM on the validation rows

0.7725

LightGBM, validation ROC AUC

0.7190

Course 4's logistic regression

0.9220

LightGBM on its own training rows

Average precision: LightGBM 0.5434, logistic regression 0.4946

```
LGBMClassifier(n_estimators=300, learning_rate=0.05, random_state=SEED, n_jobs=1, verbose=-1)
```

On rows it was fitted on, a model looks sharper than it is. Remember that for the stack.

Where the name comes from

Ke, Meng, Finley, Wang, Chen, Ma, Ye, Liu, 2017


"LightGBM: A Highly Efficient Gradient Boosting Decision Tree"

Advances in Neural Information Processing Systems 30 (NIPS 2017)

The same boosting as day 7, with trees grown and split in a faster way.

Source, read 5 October 2026: the paper's page in the NeurIPS proceedings (papers.nips.cc). Title, authors, volume, and year only; the paper itself was not read for this course.

num_leaves and min_child_samples, inside the training rows



Five folds

Defaults: 0.7745 (sd 0.0020). Fewer leaves score higher; the leaf minimum stays within a bar.

Read every difference against the fold standard deviation. Nothing is chosen today.

HistGradientBoosting: scikit-learn's own

255

bins (max_bins)

31

leaves (max_leaf_nodes)

55

trees fitted of 100

0.7773

validation ROC AUC

early_stopping="auto"

On above 10,000 training rows: it sets aside 10 percent of the rows it is given (validation_fraction), picked with random_state=SEED, and stops when they stop improving.

The rows it stops on come from the training rows: day 8's rule, kept by the library.

Three boosters, one validation set

Model	Validation ROC AUC	Average precision
logistic regression	0.7190	0.4946
LightGBM	0.7725	0.5434
HistGradientBoosting	0.7773	0.5439
XGBoost	0.7794	0.5447

0.0069

spread among the boosters

Different trees, settings, and code, within a few thousandths. Day 10 measures the noise.

Stacking on out-of-fold outputs

1. Five folds

Fit each base model on four fifths of the training rows

2. Out of fold

Keep its output on the fifth it did not see

3. Two columns

logistic and lightgbm, correlation 0.8163

4. Final model

A logistic regression on the two columns

0.7743

LightGBM out of fold

0.9220

LightGBM in sample

4.0365

coefficient, lightgbm

1.1551

coefficient, logistic

```
inputs = ensemblekit.stack_inputs(base_models, X_train, y_train)
```

StackingClassifier does the same

```
StackingClassifier(list(base_models.items()), final_estimator=LogisticRegression(),  
                  stack_method="predict_proba", cv=mlkit.cv_folds(), n_jobs=1)
```

Final model's coefficients

equal within 1e-10: True

Validation outputs


equal within 1e-12: True

cv= is always a seeded splitter, never a bare number. The base models are then refitted on all the training rows.

Does the stack beat its best member?

Model	ROC AUC	Avg precision
logistic regression	0.7190	0.4946
LightGBM	0.7725	0.5434
average of the two	0.7641	0.5354
stack (logistic layer)	0.7718	0.5454

Stack minus LightGBM: ROC AUC -0.0007, average precision +0.0020.



Not on ROC AUC (-0.0007); on average precision, yes (+0.0020). Both smaller than the boosters' spread of 0.0069.

Excel side by side

Excel side by side (tested in Excel)

`=AVERAGE(B2,C2)`

`=$H$6*B2+(1-$H$6)*C2`

`=1/(1+EXP(-(H$2+H$3*B2+H$4*C2)))`

`=AVERAGE(E2:E6001)`

the plain blend: 0.2409 on the first account

a weighted blend, weight 0.3 in H6

the stack's output: 0.2034 on the first account

average of the stack's outputs: 0.2190

No LightGBM, no boosting, no cross-validation in Excel. Intercept -2.5861 and coefficients come from Python; Excel applies them.

Save the day: the log and git blame


```
model,rows,roc_auc,average_precision,threshold,accuracy,precision,recall,f1
lightgbm,validation,0.7725,0.5434,0.1778,0.7060,0.4043,0.6956,0.5114
xgboost,validation,0.7794,0.5447,0.1812,0.7133,0.4115,0.6888,0.5152
blend_logistic_lightgbm,validation,0.7641,0.5354,0.1853,0.6987,0.3976,0.7038,0.5082
stack_logistic_lightgbm,validation,0.7718,0.5454,0.1618,0.7110,0.4103,0.7016,0.5178
```

```
git blame -s --root -L "/def stack_inputs/,+1" ensemblekit.py
```

Thresholds from out-of-fold probabilities on the training rows. All-zero baseline 0.7788, never to choose. 35 tests pass.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: the stack's columns are in-sample outputs: each base model scores the rows it was fitted on.

Code that runs is not the same as code that is right.

What you can do now

- LightGBM grows leaf-wise to num_leaves, on bins
- n_jobs=1 and verbose=-1 on every LightGBM
- Settings compared by folds of the training rows
- A stack's columns are out-of-fold outputs
- Say plainly when a stack does not beat its best member

Week 2, Day 5 Case study: boosting on the card file, and the noise in a gap