



Python for Fixed Income

COURSE 6 · NEURAL NETWORKS

Tensors and Gradients

WEEK 1 · DAY 1

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

One idea, repeated

Week 1: One neuron to a network

Gradients, layers, batches, a case study

Week 2: Training

Optimizers, dropout, class weights

Week 3: Sequences

Walk-forward on the Treasury curve

The loop inside every fit

1. Compute the loss on the rows
2. Work out the gradient of each weight
- 3. Move each weight a little downhill**
4. Repeat

Today the loop is done by hand, for one neuron, and checked against TensorFlow.

Where Courses 4 and 5 left the card file

0.7190

Course 4 logistic regression

0.7725

Course 5 LightGBM

Validation ROC AUC; average precision 0.4946 and 0.5434. 18,000 training rows fit, 6,000 validation rows report.

Day 5

A network beside both, on these validation rows, with intervals. The course says what the numbers support and promises nothing.

No network yet today: today is how a network is fitted.

One rule for the series: test rows stay closed

18,000 training rows

6,000 validation

6,000 test

fit, and from day 2 stop or choose

report only

closed

Why closed

Course 4 opened them once, on day 10. Scored again after their numbers were seen, they would become rows the course chooses against.

From day 2: fit and stop

14,400 rows fit; a stratified fifth, 3,600 rows, stops or chooses (Course 5's `early_stopping_rows`). Validation reports once per chosen model.

`course4_card_rows` never returns a closed row. Course 6 chooses nothing on the validation rows.

Devices: TensorFlow on the CPU

Devices TensorFlow can use here

CPU: `/physical_device:CPU:0`

GPUs: `0`

TensorFlow's own notice on Windows

"TensorFlow GPU support is not available on native Windows for TensorFlow >= 2.11."

Colab: Runtime, Change Runtime Type, Hardware Accelerator

Colab's FAQ: "free of charge access to computing resources, including GPUs and TPUs"; "not guaranteed and not unlimited". Not run or timed for this course.

Every notebook runs on the CPU within its time budget. A GPU is never required.

Sources, read 6 October 2026: [tensorflow/python/framework/config.py](#) (TensorFlow 2.21.0); [research.google.com/colaboratory/faq.html](#)

Tensors: a shape and a dtype

Tensor	Shape	dtype
tf.constant(0.5)	()	float32
PAY_0, training rows	(18000,)	float64
five card rows	(5, 19)	int64

Two precisions

float32: about 7 digits, Keras's default; compared within 1e-6. float64: about 16 digits; compared within 1e-12.

```
one_number * pay_0                                # InvalidArgumentError: float32 times float64
tf.cast(one_number, tf.float64) * pay_0 - 1.0      # runs, element by element
```

Rows first, then columns: a batch of rows is a range of cells.

One neuron by hand

$$z = b + w * x$$
$$p = 1 / (1 + \exp(-z))$$

One input, one weight, an intercept, the sigmoid: Course 4's logistic regression on one column.

Start: $w = 0.5$, $b = -1.0$, on PAY_0

Mean output 0.2795 against a default rate of 0.2212: too high.

0.4876

mean log loss at the start

0.471993

scikit-learn's minimum on PAY_0

Log loss: $-(y \ln p + (1 - y) \ln(1 - p))$, averaged. Keras names it `binary_crossentropy`.

The gradient by hand, and by TensorFlow

```
gradient for w = average of (p - y) * x  
gradient for b = average of (p - y)
```

-0.033916

weight: negative, so raise w

0.058284

intercept: positive, so lower b

tf.GradientTape finds the same two numbers from the loss alone: within 1e-12 in float64, 1e-6 in float32.

For one neuron the formula is two lines. For thousands of weights, the tape works it out.

Gradient descent by hand

```
w = w - learning_rate * gradient    # every weight, then again
```


Step	Weight	Mean log loss
0	0.500000	0.487621
1	0.533916	0.483442
10	0.691431	0.472997
50	0.769413	0.471994
100	0.770210	0.471993

100 steps at learning rate 1, against scikit-learn

By hand: weight 0.770210, intercept -1.418545. Scikit-learn, unpenalised: 0.770212 and -1.418548. Equal within 1e-5; loss 0.471993 both.

One neuron fitted by gradient descent is a logistic regression: day 2.

A learning rate that is too large



Learning rate 20: 0.4876, then 0.5042 after one step; as high as 1.4305; 0.7805 after 100. Too large never arrives.

What Keras automates

```
keras.utils.set_random_seed(SEED)
neuron = keras.Sequential([keras.Input(shape=(1,)), keras.layers.Dense(1, activation="sigmoid")])
neuron.compile(optimizer=keras.optimizers.SGD(learning_rate=1.0), loss="binary_crossentropy")
neuron.fit(X, y, epochs=1, batch_size=len(X), shuffle=False, verbose=0)  # one step
```

After one step	Weight	Intercept
Keras's fit	0.533916	-1.058284
gradient_step	0.533916	-1.058284

What Keras adds, days 2 to 5

A gradient for every weight, mini-batches, epochs, better optimizers, stopping rows.

Equal within 1e-6: the step you wrote by hand is the step inside fit.

The name for what Keras does to every weight

Rumelhart, Hinton, and Williams, October 1986

"Learning representations by back-propagating errors"

Nature, volume 323, issue 6088, pages 533 to 536

Back-propagation: the gradient for every weight, from the output back to the inputs.

Source, read 6 October 2026: the Crossref record of the paper (DOI 10.1038/323533a0). Bibliographic facts only; the paper itself was not read for this course.

Excel side by side

Excel side by side (tested in Excel)

`=1/(1+EXP(-($H$1+$H$2*A2)))`

`=-(B2*LN(C2)+(1-B2)*LN(1-C2))`

`=SUMPRODUCT(C2:C18001-B2:B18001)/ROWS(A2:A18001)`

`=SUMPRODUCT(C2:C18001-B2:B18001,A2:A18001)/ROWS(A2:A18001)`

`=$H$1-$H$5*$H$6` and `=$H$2-$H$5*$H$7`

`=B2-$I$1*E2` and `=C2-$I$1*F2`

the neuron, in C2, filled down

filled down; `=AVERAGE(D2:D18001)` = 0.487621

intercept gradient 0.058284

weight gradient -0.033916

one step

100 rows: weight 0.770210

No automatic differentiation and no network training in Excel. Every number above equals Python's within 1e-12.

Save the day: the experiment log

1. Init

The work folder, local settings,
the guard

2. .gitignore

adds *.keras: saved networks

3. Code

mlkit, ensemblekit, netkit and
its tests

4. Results


results.csv, 4 decimals

```
model,rows,roc_auc,average_precision,threshold,accuracy,precision,recall,f1
course4_logistic_regression,validation,0.7190,0.4946,0.5000,0.8123,0.7248,0.2442,0.3653
pay_0_neuron_by_hand,validation,0.6834,0.4241,0.5000,0.8160,0.6831,0.3135,0.4298
```

Threshold 0.5 today; from day 5, out-of-fold. Accuracy beside the all-zero baseline, 0.7788, never to choose.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: the gradient is summed, not averaged: weight gradient -610.4847, loss 2.8860 after one step.

Code that runs is not the same as code that is right.

What you can do now

- Say why TensorFlow runs on the CPU on Windows
- Read a tensor's shape and dtype; cast between them
- Compute a neuron's output and mean log loss by hand
- Check a hand gradient against `tf.GradientTape`
- Fit one neuron by gradient descent: a logistic regression

Week 1, Day 2 One neuron is a logistic regression, in Keras