



Python for Fixed Income

COURSE 6 · NEURAL NETWORKS

One Neuron Is a Logistic Regression

WEEK 1 · DAY 2

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

One model, three tools

Keras

```
Dense(1, activation="sigmoid")
```

fitted by Adam

scikit-learn

```
LogisticRegression(C=np.inf)
```

no penalty

Excel

a column of formulas

fitted by Solver

From today: fit rows and stopping rows

18,000 training rows: 14,400 to fit, a stratified fifth (3,600) to stop or choose. Validation reports once, for the model chosen.

If one neuron is a logistic regression, all three must give the same weights.

A Dense layer of one unit is day 1's formula

```
keras.layers.Dense(1, activation="sigmoid", name="neuron")  
output = sigmoid(bias + x @ kernel)
```

(2, 1)

kernel: inputs by units

(1,)

bias: the intercept

3

parameters, two inputs

Weights set by hand: predict equals the formula worked in numpy, within 1e-6.

The kernel holds the coefficients, the bias the intercept: a logistic regression.

Scaling for a network: fitted on the fit rows only

128867.52

standard deviation of LIMIT_BAL

1.11

standard deviation of PAY_0

```
scaler, A_fit, A_stop, A_valid = netkit.fit_arrays(StandardScaler(), X_fit, X_stop, X_valid)
```

Scaled column	fit mean	fit sd	stopping mean	validation mean
LIMIT_BAL	0.0	1.0	0.0318	0.0061
PAY_0	0.0	1.0	0.0172	-0.0082
BILL_AMT1	0.0	1.0	0.0127	0.0292

The scaler learns from rows, so it is fitted like a model: on the fit rows, nothing else.

Compile and fit, argument by argument

```
neuron_two.compile(optimizer=keras.optimizers.Adam(learning_rate=0.1), loss="binary_crossentropy")
neuron_two.fit(X, y, epochs=300, batch_size=len(X), shuffle=False, verbose=0)
```

Argument	What it sets
optimizer	how a step is taken: Adam, a step size per weight
loss	binary_crossentropy, day 1's mean log loss
epochs	passes over the rows
batch_size	rows per gradient; all of them is full batch
shuffle	reorder each epoch; nothing to reorder in one batch
verbose=0	no progress bar

Two columns: scikit-learn's weights

Coefficient	neuron	scikit-learn
intercept	-1.102609	-1.102607
PAY_0	0.721292	0.721293
LIMIT_BAL_10k	-0.020327	-0.020327

Adam at 0.1, full batch, 300 epochs

Mean log loss 6.2950 at epoch 1, 0.4678 at epoch 300.
Equal within 1e-4.

2.0571

odds ratio, PAY_0

0.9799

odds ratio, per 10,000 of limit

Course 4 day 6's odds ratios, to 4 decimals: they describe the training rows of this file.

Solver finds the same weights

Excel side by side (tested in Excel)

`=1/(1+EXP(-(K$1+K$2*B2+K$3*D2)))`

`=A2*LN(E2)+(1-A2)*LN(1-E2)`

`=SUM(F2:F18001)`

Solver: Max K4 by changing K1:K3

E2, filled down; K1:K3 start at 0

F2, filled down

K4, the objective

GRG Nonlinear, Course 4 day 6's options

Coefficient	Solver	neuron
intercept	-1.102607	-1.102609
PAY_0	0.721293	0.721292
LIMIT_BAL_10k	-0.020327	-0.020327

Equal within 1e-4

Solver, scikit-learn, and Keras: three tools, one model.

Nineteen columns: the same minimum

0.461056

neuron: mean log loss, fit rows

0.461056

scikit-learn, same rows

Bill weight	neuron	scikit-learn
BILL_AMT1	-0.3794	-0.3794
BILL_AMT2	0.0292	0.0292
BILL_AMT3	0.1210	0.1210
BILL_AMT4	-0.0461	-0.0461
BILL_AMT5	0.1652	0.1652
BILL_AMT6	0.0178	0.0178

All 20 coefficients equal within 1e-4

Read one at a time, the bill weights say a higher bill this month goes with a lower output. Look first at how alike the six columns are.

Near-copy columns: a flat loss

0.80 to 0.95

correlations between the six bill columns


0.3640

largest change in one weight

Mean log loss	fitted weights	bill weights averaged
fit rows	0.4611	0.4622
stopping rows	0.4685	0.4679

Many different weights fit almost equally well: compare networks by loss and scores, not weights.

Mini-batches stop short



Batch 256: 57 steps an epoch, 5,700 in all. Coefficients up to 0.1827 off (intercept); furthest weights BILL_AMT5, BILL_AMT3, BILL_AMT6. Loss 0.4648 against 0.4611 on the fit rows, 0.4728 against 0.4685 on the stopping rows.

One setting changed: the batch size, all 14,400 rows to 256; chosen on the 3,600 stopping rows: full batch.

The chosen neuron, once on validation

Model	Rows fitted on	ROC AUC	Average precision
Neuron, 19 columns	14,400 fit rows	0.7180	0.4945
scikit-learn, unpenalised	14,400 fit rows	0.7180	0.4945
Course 4 logistic regression	18,000 training rows	0.7190	0.4946

Validation, scored once. These are the bar for day 3's hidden layers.

Allowed names

model probability of default in this file

flagged at threshold t

the model's score (class-weighted)

Never called

a credit decision

an approval or a decline

a risk grade or credit score

Excel side by side

Excel side by side (tested in Excel)

```
=AVERAGE(fit!A2:A14401)
=STDEV.P(fit!A2:A14401)
=(A2-stats!A$2)/stats!A$3
=1/(1+EXP(-(stats!$B$6+SUMPRODUCT(U2:AM2,stats!$A$4:$S$4))))
=EXP(K2) and =EXP(K3)
```

stats row 2, filled right: scaler.mean_
stats row 3: scaler.scale_ (population)
valid U:AM: the fit rows' statistics
valid AO: predict
odds ratios 2.0571 and 0.9799

No Keras or network training in Excel. AVERAGE and STDEV.P equal scikit-learn's; outputs equal predict within 1e-6 on 6,000 rows.

Save the day: git diff


```
<hash> Experiment log: the 19-column neuron fitted in Keras, validation metrics  
<hash> Experiment log as day 1 left it  
<hash> netkit: fit_arrays and neuron_coefficients, 10 tests
```

```
+neuron_19_columns,validation,0.7180,0.4945,0.5000,0.8115,0.7130,0.2472,0.3671
```

One + line, no row removed: 4 decimals and LF line endings keep the diff clean.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: the scaler is fitted on the fit and validation rows: means off 0 by 0.0085; LIMIT_BAL mean 166,784.12, not 166,552.78.

Code that runs is not the same as code that is right.

What you can do now

- Read a Dense layer's kernel and bias
- Scale with `fit_arrays`, on the fit rows only
- Show one neuron is scikit-learn's logistic regression
- Fit the same weights with Solver in Excel
- Read near-copy weights with care; choose on stopping rows

Week 1, Day 3 Hidden layers and activations