



Python for Fixed Income

COURSE 6 · NEURAL NETWORKS

Batch Size, Epochs, and Loss Curves

WEEK 1 · DAY 4

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

Three settings, two curves

Batch size

How many rows each step looks at

Epochs

How many passes over the fit rows

Early stopping

When to end the fitting, by a callback

Two curves, recorded every epoch

Training loss: the 14,400 fit rows

Stopping loss: the 3,600 stopping rows

Every setting judged on the stopping rows

Validation rows (6,000): report only

Last section: the seed spread, the yardstick for every comparison after today.

An epoch, a batch, a step

Batch size	Steps per epoch
32	450
64	225
256	57
1,024	15
14,400	1


Three words

A batch: the rows one step's gradient is worked out from. A step: one move of every weight. An epoch: one pass through all the batches.

```
netkit.steps_per_epoch(14_400, 64) # ceil
```

At the same epochs, a smaller batch takes more steps, each from fewer rows.


Two loss curves, one network



Training loss falls every epoch (0.4930 to 0.4055); stopping loss: 0.4700, lowest 0.4459, then 0.4550.

Batch size at a fixed 20 epochs

Batch	Steps	Train loss	Stop loss	Lowest stop	At epoch
32	9000	0.4134	0.4478	0.4440	5
64	4500	0.4149	0.4486	0.4459	8
256	1140	0.4187	0.4434	0.4432	17



Compare each at its own best epoch

At one fixed epoch the batches sit at different points of their curves. Their lowest stopping losses are close: section 4.8 says how close is close.

Smaller batches fit more per epoch and turn sooner; larger ones turn later.

The one cell outside the size

```
# thread tolerance: large batches differ in the last digits across thread counts  
batch 1,024: 300 steps in all; training loss 0.430, stopping loss 0.446  
batch 14,400: 20 steps in all; training loss 0.534, stopping loss 0.532
```

Why the last digits move

A large batch is a large product. On every core TensorFlow adds its parts in a different order, and float32 sums differ in the last digits.

The rule for this one cell

Marked on its first line, printed to 3 decimals, allowed to move by up to 0.005. Every other cell matches exactly; nothing later uses it.

Larger batches take fewer steps per epoch, so at the same epochs they move the weights less.

The best epoch is read on the stopping loss

8

best epoch on the stopping loss, of 40

40

the "best" epoch on the training loss: the last

	Epoch 8	Epoch 40
Stopping loss	0.4459	0.4550
Training loss	0.4240	0.4055

`netkit.best_epoch`

`np.argmin` of `val_loss`, plus 1: the first epoch with the smallest stopping loss.

The training loss keeps falling for as long as you fit: it cannot pick the epoch.

Early stopping: the epoch it picked, the weights it kept

```
keras.callbacks.EarlyStopping(monitor="val_loss", patience=10, restore_best_weights=True)
model.fit(..., epochs=200, batch_size=64, callbacks=[stopper], verbose=0)
```

18

epochs run, of up to 200

8

epoch kept (EarlyStopping's best_epoch + 1: 8)

0.4459

stopping loss of the weights kept (smallest:
0.4459)

The same run as section 4.2, epoch for epoch, until it stopped: early stopping changed when the fitting ended.

Validation: the same ranking, a better log loss


Validation rows	ROC AUC	Average precision	Mean log loss
64-32, 40 epochs	0.7689	0.5291	0.4460
64-32, early stopping	0.7689	0.5258	0.4430

Where the extra epochs show

On the log loss: the network fitted past its best epoch gives outputs more confident than the rows support, while the order of the accounts hardly changes.

Reported once each. The epoch was chosen on the stopping rows, never on these.

Reading a curve: underfit, fit, overfit



The neuron ends at 0.4621 and 0.4702, still falling; its lowest stopping loss 0.4702 is above the network's 0.4459.

Where a GPU would help

Large products

A step multiplies each batch by each layer's weights. A GPU does many multiplications at once, so it helps with large batches through wide layers.

This course's networks

On the card file: 19 inputs, hidden layers of 64 units or fewer, batches of 256 or fewer outside one marked cell. Small products: the CPU is enough.

Colab: Runtime, Change Runtime Type. Never required; nothing here was run or timed on a GPU.

Day 1's devices paragraph is repeated word for word in the notebook, with its sources (TensorFlow 2.21.0's config.py; Colab's FAQ, read 6 October 2026).

Seeds: the yardstick for every comparison

Seed	Best epoch	Lowest stop loss	Validation ROC AUC
SEED + 0	8	0.4459	0.7689
SEED + 1	13	0.4471	0.7697
SEED + 2	12	0.4440	0.7701
SEED + 3	11	0.4486	0.7696
SEED + 4	8	0.4462	0.7712

0.0046

range of the lowest
stopping loss (sd
0.0015)

0.0023

range of validation ROC
AUC (sd 0.0007)

Section 4.3's batches

Lowest stopping losses within 0.0027: inside the
seed's range.

Five validation scores measure the spread; picking a seed on them would be choosing on the report rows.

Excel side by side

Excel side by side (tested in Excel)

`=ROUNDUP(14400/64,0)`

`=ROUNDUP(14400/F2,0)`

`=MATCH(MIN(C2:C41),C2:C41,0)`

`=MIN(C2:C41)`

`=MATCH(MIN(B2:B41),B2:B41,0)`

`=C2-B2`

`=MATCH(MIN(K2:K5),K2:K5,0)`

225 steps per epoch

filled down: 450, 225, 113, 57, 15, 1

epoch 8, the best epoch

0.445864, the lowest stopping loss

epoch 40: the AI check's mistake

filled down: the gap, exercise 3's curve_gap

2 on 0.7, 0.5, 0.6, 0.5: the first of a tie

No network training in Excel: it cannot make a loss curve, but it reads one. Equal to Python within 1e-12.

Save the day: git show HEAD~1

1. Modules

mlkit, ensemblekit, netkit and its tests

2. 40 epochs

results.csv, the first row

3. Early stopping


results.csv, a second row

```
!git -C "{work}" show HEAD~1:results.csv
model,rows,roc_auc,average_precision,threshold,accuracy,precision,recall,f1
dense_64_32_40_epochs,validation,0.7689,0.5291,0.5000,0.8193,0.6431,0.4115,0.5018
```

HEAD~1:results.csv is the file one commit back, folder unchanged. All-zero baseline accuracy 0.7788.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: the epoch is read from the training loss: it returns 18, the last epoch run, where EarlyStopping kept 8.

Code that runs is not the same as code that is right.

What you can do now

- Count the steps in an epoch at any batch size
- Draw and read the training and stopping loss curves
- Judge each batch size at its own best epoch
- Find the best epoch on the stopping loss; stop early
- Read a comparison against the seed spread

Week 1, Day 5 Case study: a first network on credit card default