



Python for Fixed Income

COURSE 6 · NEURAL NETWORKS

Momentum, Adam, and Learning Rate Schedules

WEEK 2 · DAY 2

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

The optimizer is the rule for the step

The loop inside every fit

1. Compute the loss on a batch
2. Work out each weight's gradient
- 3. Move each weight: the optimizer**
4. Repeat

Four rules for step 3

SGD	learning rate x gradient
Momentum	keeps part of the last step
Adam	gradient / its running size
RMSprop	Adam without m or correction

First on day 1's neuron, checked against Keras; then on the 64-32 network, chosen on the stopping rows.

Plain SGD, as Keras writes it

```
self.assign_sub(variable, ops.multiply(gradient, learning_rate)) # keras/src/optimizers/sgd.py
```

0.01

SGD's default learning rate

0

SGD's default momentum

One step, learning rate 1	Intercept	Weight
Keras SGD, by apply	-1.058284	0.533916
netkit.gradient_step	-1.058284	0.533916


optimizer.apply(gradients, variables): the optimizer on its own, as fit uses it.

Equal within 1e-6: Keras's plain step is day 1's gradient_step.

Momentum by hand: it carries the last step

```
velocity = momentum * velocity - lr * gradient
weight   = weight + velocity
```

Two steps	Intercept	Weight
momentum_step	-1.159484	0.593217
Keras SGD(0.9)	-1.159484	0.593217



Learning rate 1: plain steps settle from step 20, momentum 0.9 from step 48. Momentum overshoots.

Adam by hand, and the bias correction

```
m = m + (g - m) * (1 - beta_1);    v = v + (g**2 - v) * (1 - beta_2)
alpha = lr * sqrt(1 - beta_2**t) / (1 - beta_1**t)
w = w - alpha * m / (sqrt(v) + epsilon)
```

First step moves	Intercept	Weight
Plain step, rate 1	0.058284	0.033916
Adam, rate 0.05	0.049997	0.049995
Adam without $1 - \beta_1^t$	0.005000	0.005000
Adam with no correction	0.158105	0.158099

Alpha at steps 1, 2, 3

0.015811, 0.011766, 0.010101

The correction fades toward the learning rate. Three steps equal Keras's Adam within 1e-6.

Adam moves every weight by the learning rate at first, whatever its gradient.

Where the name Adam comes from

Diederik P. Kingma and Jimmy Ba, 22 December 2014

"Adam: A Method for Stochastic Optimization"

arXiv 1412.6980; published at the 3rd International Conference for Learning Representations, 2015

`keras.optimizers.Adam` carries its name from this title.

Source, read 6 October 2026: the arXiv record of the paper, arxiv.org/abs/1412.6980. Bibliographic facts and the record's comment only; the paper itself was not read for this course.

RMSprop, and the defaults Keras reads out

Optimizer	learning_rate	momentum	rho	beta_1	beta_2	epsilon
SGD	0.01	0	-	-	-	-
RMSprop	0.001	0	0.9	-	-	1e-07
Adam	0.001	-	-	0.9	0.999	1e-07


```
v = rho * v + (1 - rho) * g**2
w = w - lr * g / sqrt(v + epsilon)      # epsilon inside the square root
```

One step at 0.05 on the neuron: both weights move 0.1580, about 3.16 times the rate.

From `get_config`: SGD defaults to 0.01, RMSprop and Adam to 0.001.


Four optimizers on one network, 20 epochs

Optimizer	Best epoch	Best stopping loss	Training loss, epoch 20
SGD, learning rate 0.01	20	0.4473	0.4356
SGD, momentum 0.9	11	0.4437	0.4188
RMSprop, learning rate 0.001	9	0.4414	0.4168
Adam, learning rate 0.001	8	0.4459	0.4149



The lowest training loss is not the lowest stopping loss. The stopping rows choose.

Adam's learning rate: too small, too large



Learning rate	Best	Epoch
0.0001	0.4472	20
0.001	0.4459	8
0.01	0.4406	7
0.1	0.4719	13

0.01 is lowest but ragged: largest rise between epochs 0.0040.

Too small is slow; too large never arrives (day 1's neuron at learning rate 20 did the same).

A schedule: start large, then halve

$$\text{lr}(t) = 0.01 * 0.5 ** (t / 1125)$$

Epoch	Step	Learning rate
1	0	0.010000
6	1,125	0.005000
11	2,250	0.002500
16	3,375	0.001250
20	4,275	0.000718

Against a constant 0.01

4,500 steps; learning rate now 0.000625. Best stopping loss 0.4418 at epoch 10, not lower. Largest rise between epochs 0.0011, against 0.0040.

Calmer later epochs; no better fit on this network. ReduceLROnPlateau is the other common choice.

The choice, then the validation rows once

Lowest stopping loss of 8 settings: Adam, learning rate 0.01, epoch 7

Validation rows, once	Best epoch	ROC AUC	Average precision
Adam 0.01 (chosen)	7	0.7745	0.5291
Adam 0.001 (default network)	8	0.7689	0.5258

cv_network, 5 folds: average precision 0.5369 (sd 0.0074); threshold for recall 0.70 0.1867

Small gaps: day 5's bootstrap intervals are the tool. Today reports the numbers and claims nothing beyond them.

Excel side by side

Excel side by side (tested in Excel)

=K\$2*G2-K\$1*E2 then =B2+G3

=G2+(E2-G2)*(1-\$N\$2)

=I2+(E2^2-I2)*(1-\$N\$3)

=\$N\$1*SQRT(1-\$N\$3^A3)/(1-\$N\$2^A3)

=B2-K3*G3/(SQRT(I3)+\$N\$4)

=\$B\$1*\$B\$2^(A6/\$B\$3)

momentum, 100 rows: weight 0.771070

Adam's m

Adam's v

alpha 0.015811 at step 1

Adam, 3 rows: weight 0.646084

0.000625 at step 4,500

No optimizer and no network training in Excel. Every number equals Python's within 1e-12 (foundation: =\$B\$1*C2-\$B\$2*D2 on hand values too).

Save the day: git stash an experiment

1. Commit

Four modules, today's two functions

2. Log

results.csv and optimizers.csv

3. Stash

push -m: the experiment off the folder

4. Drop


show -p to compare, then drop

```
model,rows,roc_auc,average_precision,threshold,accuracy,precision,recall,f1
adam_0_01_64_32,validation,0.7745,0.5291,0.1867,0.6453,0.3584,0.7641,0.4880
```

The experiment: SGD without momentum at 0.1. Shelved, compared, dropped; git stash pop would bring it back.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: half the bias correction is missing: the first step moves each weight 0.0050 instead of 0.05.

Code that runs is not the same as code that is right.

What you can do now

- Check an optimizer against Keras with apply
- Take momentum steps; say why momentum overshoots
- Take Adam steps; say what the bias correction does
- Choose optimizer and learning rate on stopping rows
- Shelf an experiment with git stash

Week 2, Day 3 Initialization and batch normalization