



Python for Fixed Income

COURSE 6 · NEURAL NETWORKS

Initialization and Batch Normalization

WEEK 2 · DAY 3

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

Before and between the steps

Where the weights start

Glorot uniform, Keras's default

He normal, for ReLU

Why not zeros

Vanishing gradients in six sigmoid layers

Between the layers

Batch normalization by hand

Four deep networks, one change each

Moving averages at scoring time

Training mode as a bug

Card file: 14,400 rows fit, 3,600 rows stop or choose, 6,000 validation rows report once.

What a layer starts from

Glorot uniform (Keras's default)

```
limit = sqrt(6 / (fan_in + fan_out))
```

19 inputs, 64 units: limit 0.2689. A fresh layer's weights run from -0.2687 to 0.2687; biases start at 0.

He normal (paired with ReLU)

```
standard deviation = sqrt(2 / fan_in)
```

3.8 million draws: 0.32445 against the formula's 0.32444.

Cut at 2 standard deviations after widening by 1 / 0.8796: largest draw 0.7377, 2.27 sd.

Source: `keras/src/initializers/random_initializers.py`, Keras 3.15.1, read 6 October 2026.

A start of zeros: sixteen copies of one unit


Start	Different hidden units	Stopping loss after 5 epochs
zeros	1	0.4706
Glorot uniform	16	0.4646

Why

Equal weights give equal outputs, equal gradients, equal steps: the units stay equal forever. A random start breaks the tie; the seed makes it the same on every run.

16 sigmoid units, Adam at 1e-3, batch 64, 5 epochs, on the fit rows; loss on the stopping rows.

Vanishing gradients in six sigmoid layers



Output over first layer: sigmoid 26,544 times (1.22×10^0 against 4.60×10^{-5}); ReLU with He 5.3 times.

Batch normalization by hand

$$(x - \text{batch mean}) / \sqrt{\text{batch population variance} + 0.001} * \gamma + \beta$$

0.4725 to 0.5473

column means before

0.9449 to 0.9898


standard deviations after

batch_norm_by_hand equals keras.layers.BatchNormalization in training mode within 1e-5.

A real batch: the first hidden layer's 32 sigmoid outputs for 64 fit rows (standard deviations 0.0912 to 0.2196 before).

Four deep networks, one setting at a time

Model	Epochs run	Best epoch	Best stopping loss
model_1_sigmoid	40	40	0.4528
model_2_sigmoid_batch_norm	21	11	0.4379
model_3_relu_he	17	7	0.4466
model_4_relu_he_batch_norm	12	2	0.4505



Lowest stopping loss: model_2_sigmoid_batch_norm. Gaps of a few thousandths: day 4's seed spread is the yardstick.

At scoring time: the moving averages

Per unit	Updated by
gamma, beta	the optimizer
moving mean, moving variance	averaging (momentum 0.99)

Model 2

6,721 numbers: 6,337 learned, 384 moving averages.

Scoring mode (predict): account 0 alone 0.2689, among the 6,000 0.2689

Accounts 0, 1, 2 in scoring mode: 0.2689, 0.0481, 0.1756

Training mode, one at a time: 0.167, 0.167, 0.167. One row: no spread, only beta.

And each training-mode run moved the moving averages: the network changed by being scored.

Score with `model.predict`, always. `training=True` only in a cell labelled demonstration.

The chosen network, scored once

0.7704

validation ROC AUC

0.5283

validation average precision

Threshold for recall 0.70, from the stopping rows: 0.1949

On the validation rows: recall 0.6797, precision 0.4020, 2,244 accounts flagged. A threshold carries over only approximately.

Model probability of default in this file. Labelled validation, never test.

The name of the layer

Sergey Ioffe and Christian Szegedy, 11 February 2015

"Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift"

arXiv 1502.03167, first version

What the layer did on the card file, you measured in section 8.5.

Source, read 6 October 2026: the arXiv record of the paper, arxiv.org/abs/1502.03167. Bibliographic facts only; the paper itself was not read for this course.

Excel side by side

Excel side by side (tested in Excel)

`=SQRT(6/(19+64))`

`=SQRT(2/19)`

`=(A2-AVERAGE(A$2:A$65))/SQRT(VAR.P(A$2:A$65)+0.001)`

`=AVERAGE(AH2:AH65)`

`=VAR.P(AH2:AH65)`

Glorot limit 0.268866

He standard deviation 0.324443

in AH2, across to BM, down to 65

0 within 1e-12

$v / (v + 0.001)$: just under 1

All 2048 cells equal batch_norm_by_hand within 1e-12 and the Keras layer within 1e-5. VAR.P, not VAR.S: the layer divides by n.

Save the day: git diff --stat HEAD~1

1. Morning

netkit as day 7 left it

2. Code

two functions, four tests

3. Results

results.csv, one row


```
netkit.py | 45 ++++++
test_netkit.py | 38 ++++++
2 files changed, 83 insertions(+)
```

```
model,rows,roc_auc,average_precision,threshold,accuracy,precision,recall,f1
model_2_sigmoid_batch_norm,validation,0.7704,0.5283,0.1949,0.7055,0.4020,0.6797,0.5052
```

Accuracy beside the all-zero baseline, 0.7788, never to choose.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: `model(A, training=True)`: ROC AUC 0.7659 looks fine, but account 0 is 0.1774 alone and 0.1921 among 6,000.

Code that runs is not the same as code that is right.

What you can do now

- Say what a Dense layer starts from, and what He normal changes
- Explain why a start of zeros fails
- Measure vanishing gradients with `gradient_norms`
- Do batch normalization by hand, in Python and Excel
- Score in scoring mode: moving averages, never `training=True`

Week 2, Day 4 Overfitting, dropout, and early stopping