



Python for Fixed Income

COURSE 6 · NEURAL NETWORKS

Overfitting, Dropout, and Early Stopping

WEEK 2 · DAY 4

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

Overfitting, and three brakes

Course 4's Polish training rows

4,680

companies

326

bankrupt

112 network inputs into a 64-32 network: thousands of weights for 326 bankruptcies.

Three brakes, each measured

L2: a price on large weights


Dropout: a fifth of values off, in training

Early stopping: keep the best epoch

Then all three by five folds

Few rows and many weights: a network can fit its own rows, noise and all.

Overfitting made visible



Stopping loss 0.1169 at epoch 13, 0.2809 at epoch 60; fit loss 0.0559. Average precision after 60: 0.7109.

L2: a price on large weights

```
loss = mean log loss + 0.001 * (sum of squared hidden kernel weights)
keras.layers.Dense(64, activation="relu", kernel_regularizer=keras.regularizers.L2(0.001))
```

Keras's stopping loss 0.1597 = log loss 0.1274 + penalty 0.0323: the penalty is in every loss it reports.

After 60 epochs	Stopping log loss	Average precision	ROC AUC	Squared weights
plain	0.2604	0.7109	0.9049	203.1065
L2 0.001	0.1274	0.7865	0.9421	32.2599

Compare an L2 network by the log loss of its outputs, never by its val_loss.

Dropout by hand

```
dropped = values * (uniforms >= 0.2) / (1 - 0.2) # netkit.dropout_by_hand
```

20 values, seeded uniforms


3 values whose draw fell below 0.2 become 0; the rest are multiplied by 1.25, so each value keeps its average.

Training against scoring

In fit: a fresh fifth dropped at every batch. In predict: nothing dropped, nothing divided; the layer passes values through.

Dropout happens while a network is fitted, never while it is scored.

Dropout against no brake



After 60 epochs: stopping log loss 0.1339 against 0.2604; average precision 0.7829 against 0.7109.

The paper that named the layer

Srivastava, Hinton, Krizhevsky, Sutskever, and Salakhutdinov, 2014

"Dropout: A Simple Way to Prevent Neural Networks from Overfitting"

Journal of Machine Learning Research, volume 15, number 56, pages 1929 to 1958

Its title holds both of today's words: the layer, and the problem it is for.

Source, read 6 October 2026: the paper's page at the journal, jmlr.org/papers/v15/srivastava14a.html. Bibliographic facts only; the paper itself was not read for this course.

Early stopping, with the best weights kept

```
stop = keras.callbacks.EarlyStopping(monitor="val_loss", patience=10, restore_best_weights=True)
model.fit(A_fit, y_fit, validation_data=(A_stop, y_stop), epochs=200, callbacks=[stop], verbose=0)
```

23

epochs run of 200

13

best epoch, weights kept

0.7535

average precision (60 epochs: 0.7109)

The stopping rows chose the epoch, so their score flatters it a little: section 9.5 scores by folds.

Same seed, same rows: the callback kept the epoch the chart marked.

The four together, by five folds

Model, five folds inside 4,680 rows	Average precision	Fold sd	ROC AUC
Course 4 logistic regression	0.6624	0.0780	0.9281
Network, early stopping	0.7353	0.0380	0.9375
Network, L2 + dropout + stopping	0.7559	0.0387	0.9425
LightGBM defaults	0.7783	0.0469	0.9509

The brakes gain less than a fold sd; LightGBM stays ahead by less than a fold sd. Day 10 adds intervals.

How many rows a network needs

Fit rows	Bankrupt	Inputs	Best epoch	Lowest stopping loss	Average precision
3,744	261	112	13	0.1169	0.7535
1,872	131	95	8	0.1413	0.6777

Both stopped and read on the same 936 stopping rows. The indicator columns follow the rows the transformer is fitted on, so the half network has 95 inputs.

The rows a network has are the first limit on what it can fit, before any setting.

What the network is not

Allowed names

model probability of bankruptcy in this file

flagged at threshold t

the model's score (class-weighted)

Never called

a credit view on any company

a rating

an investment decision

Allowed names

weight

gradient

the network leans on

Never called

driver or key driver

cause or reason

what the network knows

A brake changes how the network fits its rows. It says nothing about why a company failed.

Excel side by side

Excel side by side (tested in Excel)

`=A2*(B2>=0.2)/(1-0.2)`

`=COUNTIF(C2:C21,0)`

`=0.001*SUMSQ(A1:BL112,BN1:CS64)`

`=MATCH(MIN(B2:B61),B2:B61,0)`

filled down: dropout_by_hand within 1e-12

3 values dropped

0.032260, l2_penalty's number

epoch 13, best_epoch's number

No network training and no early stopping in Excel. Each brake's arithmetic is a formula.

Save the day: a change inside a line

```
git diff --word-diff --word-diff-regex="[,]+" results.csv
```


```
network_l2_dropout_early_stopping,cv,0.9425,0.7559,[-0.5000,0.9626,0.8578,0.5552,0.6741-]  
{+0.2006,0.9489,0.6173,0.7025,0.6571+}
```

[-old-] and {+new+} inside the line: only the threshold and the four flag columns change. ROC AUC and average precision stay, because a threshold does not reorder the companies.

A word is anything between commas: in a CSV, a cell.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: scored with training=True: average precision 0.7530, then 0.7106, against predict's 0.7901.

Code that runs is not the same as code that is right.

What you can do now

- Read loss curves and say where overfitting starts
- Price large weights with L2, and allow for the penalty
- Do inverted dropout by hand; score with predict
- Stop early and keep the best epoch's weights
- Compare brakes by five folds beside LightGBM

Week 2, Day 5 Case study: class weights on bankruptcy