



Python for Fixed Income

COURSE 6 · NEURAL NETWORKS

Recurrent Networks: LSTM and GRU

WEEK 3 · DAY 2

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

Where day 11 left the curve

The question, in the past tense

For each date in the development years: the 10 Yr change in bp, from a model fitted only on earlier dates, using the 20 changes before it.

3.21

zero change (bp)

3.22

training mean (bp)

3.35

day 11's dense network (bp)

Today: layers that read the window in order, one day at a time, and carry a state.

A walk-forward error on past dates is not evidence about any later date, and not a view on rates.

What the output may be called

Allowed names

fitted change in bp

walk-forward error

the last value carried forward

Never called

a forecast or a prediction

tomorrow's move or a call

an outlook or a view on rates

Development years: 2,001 dates, 2015-01-02 to 2022-12-30. The held-out years open once, on day 14.

A recurrent step, carried along the window

```
h = zeros(16)                                # before day 1
h = tanh(x[t] * W + h @ U + b)                # for each day t of the window
```

Weight	Shape	What it does
W, the kernel	(1, 16)	today's input, one weight per unit
U, the recurrent kernel	(16, 16)	yesterday's state into today's
b, the bias	(16,)	an intercept per unit

Against Keras

Day 20 state within 1e-6:
True

First window 2015-01-05 to 2015-02-02, scaled with fold 1's scaler (mean 0.0707 bp, sd 4.9244 bp); its target 2015-02-03 is not in it.

Counting a recurrent layer's weights

```
SimpleRNN  u (u + f + 1)
LSTM       4 u (u + f + 1)
GRU        3 u (u + f + 2)
```

Layer, 16 units	1 channel	11 channels
SimpleRNN	288	448
LSTM	1,152	1,792
GRU	912	1,392

The window's length is in no formula

The same weights are used on every day, so 20 days and 10 days need the same count. Day 11's dense layer had a weight per day.

Formula equal to count_params in every row; 53 tests pass by the end of the day.

The LSTM's gates

Input gate

sigmoid: how much to write

unit 1, day 1: 0.4118

Forget gate

sigmoid: how much to keep

unit 1, day 1: 0.8145

Candidate

tanh: what today would write

unit 1, day 1: 0.3701

Output gate

sigmoid: how much to show

unit 1, day 1: 0.5928

```
c = forget * c + input * candidate
h = output * tanh(c)
```

Weights four times as wide: (1, 64), (16, 64), (64,). Forget intercepts start at 1.

lstm_step carried 20 days equals the Keras layer within 1e-6: True.

The name on the layer

Sepp Hochreiter and Juergen Schmidhuber, 1 November 1997

"Long Short-Term Memory"

Neural Computation, volume 9, issue 8, pages 1735 to 1780

`keras.layers.LSTM`: the cell state a step can keep or forget through its gates.

Source, read 6 October 2026: the Crossref record of the paper (DOI 10.1162/neco.1997.9.8.1735). Bibliographic facts only; the paper itself was not read for this course.

The GRU: one state, two gates

Update gate

Mixes yesterday's state with today's candidate.

Reset gate

Sets how much of yesterday's state goes into the candidate.

Kernel (1, 48), recurrent kernel (16, 48), bias (2, 48): 912 weights on one channel.

Two bias rows (`reset_after=True`, Keras's default): the + 2 in $3u$ ($u + f + 2$).


Usually traced to Cho and six co-authors, arXiv 1406.1078, first version 3 June 2014 (the arXiv record, read 6 October 2026; the paper was not read for this course).

Walk-forward: three recurrent networks

Rule or network	Mean fold MAE (bp)	Folds below zero change
zero change (last value carried forward)	3.21	
training mean	3.22	0 of 5
dense (day 11's, rebuilt)	3.35	0 of 5
SimpleRNN	3.31	0 of 5
LSTM	3.23	0 of 5
GRU	3.22	1 of 5

Every network scored more than zero change. The GRU, closest, is chosen on the development years.


Fold by fold: beside zero change



Folds below zero change: dense 0 of 5, SimpleRNN 0 of 5, LSTM 0 of 5, GRU 1 of 5. Every line moves with the fold.

Why a network ends near the mean

Network	Spread ratio	Best epochs by fold
dense	0.3037	[50, 28, 25, 11, 8]
SimpleRNN	0.1947	[9, 16, 14, 8, 20]
LSTM	0.0404	[1, 5, 1, 1, 3]
GRU	0.0603	[1, 4, 2, 1, 2]



Fitted sd over actual sd (4.3254 bp). 0 would be one number on every date; best epoch 1 is a network that barely moved.

Closest to zero change by fitting little; the networks that fitted more scored worse.

One setting changed: the window

GRU	MAE (bp)	Zero change (bp)	Folds below	First test date
20-day window	3.22	3.21	1 of 5	2017-02-07
10-day window	3.20	3.19	1 of 5	2017-01-24

Each window against its own zero change

A shorter window starts its targets earlier, so its folds hold slightly different dates.

One setting changed from the GRU: the window, 20 to 10 trading days; chosen on the development years' five folds.

Excel side by side

Excel side by side (tested in Excel)

```
=TANH($B3*rw!$A$1:$P$1+MMULT(C2:R2, rw!$A$3:$P$18)+rw!$A$2:$P$2)  
=TANH($B4*rw!A1+SUMPRODUCT(C3:R3, TRANSPOSE(rw!A3:A18))+rw!A2)  
=1/(1+EXP(-AY3:BN3))*S2:AH2+1/(1+EXP(-AI3:AX3))*TANH(B03:CD3)  
=1/(1+EXP(-CE3:CT3))*TANH(S3:AH3)  
=AVERAGE(ABS(B2:B253-C2:C253))  
=STDEV.P(C2:C1261)/STDEV.P(B2:B1261)
```

SimpleRNN, 16 units, spilled; filled down 20 days
one unit, by SUMPRODUCT
LSTM cell state (forget, input, candidate)
LSTM hidden state (output gate)
GRU MAE, fold 1: 2.7576 bp
spread ratio 0.0603

No recurrent layer, training, or walk-forward in Excel. Every number equals Python's within 1e-12.

Save the day: four commits

1. Morning

netkit as day 11 left it

2. Code

two functions, five tests

3. Day 11's log

rows on changes, rebuilt

4. Today's rows


four networks added

```
model,channels>window,rows,folds,mae_bp,zero_change_mae_bp,folds_below_zero_change
zero_change,1,20,development,5,3.2087,3.2087,0
training_mean,1,20,development,5,3.2216,3.2087,0
dense_16_changes,1,20,development,5,3.3472,3.2087,0
simple_rnn_16,1,20,development,5,3.3099,3.2087,0
lstm_16,1,20,development,5,3.2281,3.2087,0
gru_16,1,20,development,5,3.2234,3.2087,1
gru_16_window_10,1,10,development,5,3.2003,3.1865,1
```

git show --stat HEAD: 1 file changed, 4 insertions(+). Added rows only.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: one scaler fitted on every development row: fold 1's mean 0.0880 bp, not 0.0707 bp.

Code that runs is not the same as code that is right.

What you can do now

- Carry a recurrent step along a window by hand
- Count a recurrent layer's weights
- Match an LSTM with lstm_step, gate by gate
- Score SimpleRNN, LSTM, GRU against zero change
- Read a spread ratio and the best epochs

Week 3, Day 3 One-dimensional convolutions and many tenors