



Python for Fixed Income

COURSE 7 · NLP AND LLMS

Sentence Embeddings and Semantic Search

WEEK 2 · DAY 1

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

From word vectors to a sentence model

Day 4: one vector per word

Looked up in a table, the same in every sentence; a document is the mean of its words.

Today: one vector per passage

Computed from all its words in their order. A question becomes a vector too.

22,713,216

parameters, MiniLM

384

numbers per text

256

word pieces kept per text

Real documents, described

A search result on the indenture: a section, a score, and its own words

Allowed names

the indenture states

extracted value, with its quote

passed the source check

Never called

loose, tight, weak, strong

covenant-lite or protective

a credit view or recommendation

A score says how close two texts sit in the model's numbers, not what the indenture means.

Sections into chunks

107

sections

454

chunks, 348 of the full 180 words

997

words before 1.01, in no chunk

```
chunks = textkit.chunk_text(sections, indenture)    # 180 words, the next 150 later
# Section 4.12: 1,819 words, 12 chunks, the last 169 words, clause (e) kept
```

Most chunks: 1.01 113, 12.14 72, 2.05 30. Section 12.14's run on through the signatures and the exhibits.

Source: Boyd Gaming Corporation, indenture for the 4.750% Senior Notes due 2031, Exhibit 4.1 to the 8-K filed 2021-06-08, accession 0001193125-21-185495.

A sentence model, and where its files went

11

files in the Hugging Face cache

91,578,415

bytes on disk

1

download, then no network

```
~\.cache\huggingface\hub      # on Windows; HF_HOME moves it
hf cache ls                  # each model and its size
hf cache rm model/sentence-transformers/all-MiniLM-L6-v2
```

The first load, download included, took about 7 seconds on the build machine.

Where sentence embeddings come from

The paper

Sentence-BERT, Reimers and Gurevych: arXiv 1908.10084, 2019-08-27, EMNLP 2019

Its abstract

"semantically meaningful sentence embeddings that can be compared using cosine-similarity"

Model card

"input text longer than 256 word pieces is truncated"

Both statements are measured on the indenture today. The paper itself was not read.

Sources, read 8 October 2026: arxiv.org/abs/1908.10084; huggingface.co/sentence-transformers/all-MiniLM-L6-v2 (pinned revision)

Embeddings, normalized, embedded once

```
matrix = textkit.embed(model, chunks["text"].tolist())  
# 454 rows of 384 numbers, 0.70 MB; chunk 0 starts -0.0977, 0.0184, -0.0135  
# dot product equal to the cosine within 1e-6 on all 206,116 pairs  
# 8 chunks embedded alone: 0 of 8 rows bit-identical, all within 1e-6
```

Embed a list once, reuse the array, compare within 1e-6, print scores at 4 decimals.

Semantic search

Question	Top section	Title	Score
Pay a dividend?	4.15	Limitation on Restricted Payments	0.5111
The company is sold?	5.01	Merger, Consolidation and Sale of Assets	0.5423
How much debt?	1.01	Definitions	0.5018
Redeem the notes early?	3.07	Optional Redemption	0.7518
Pledge assets as collateral?	1.01	Definitions	0.5540

A search ranks chunks. Reading the section says what it answers.

Source: Boyd Gaming Corporation, indenture for the 4.750% Senior Notes due 2031, Exhibit 4.1 to the 8-K filed 2021-06-08, accession 0001193125-21-185495.

Subword tokens and the 256 limit

Word	Word pieces
Indebtedness	ind, ##eb, ##ted, ##ness
EBITDA	e, ##bit, ##da
Consolidated	consolidated
Restricted	restricted

34

of 454 chunks over
256 pieces

316

pieces, the longest
chunk

140

words kept, the
fewest

Median 224 pieces a chunk; 50 of 107 whole sections over 256 (4.12 alone 2,561).

A sentence added past piece 256: identical vector (True). The model cuts silently.

The same questions by TF-IDF

Question	MiniLM	Score	TF-IDF	Score
Pay a dividend?	4.15	0.5111	4.16	0.2142
The company is sold?	5.01	0.5423	2.01	0.1876
How much debt?	1.01	0.5018	12.09	0.3299
Redeem the notes early?	3.07	0.7518	3.07	0.2542
Pledge assets as collateral?	1.01	0.5540	1.01	0.1339
Football match (control)	6.08	0.0887	1.01	0.0000

Word overlap needs the document's own words; 2,915 columns, stop words out.

A score is not a probability

0.7518

redemption, Section 3.07

0.5111

dividend, Section 4.15

0.0887

football, Section 6.08

A cosine between -1 and 1. Not a probability, not a confidence; it ranks chunks for one question and does not compare questions. top_k always returns k rows.

A search gives places to read. The reading stays with you.

Excel side by side

Excel side by side (tested in Excel)

`=SUMPRODUCT(A1:A384,B1:B384)`

`=SUMPRODUCT(A1:A384,B1:B384)/(SQRT(SUMSQ(A1:A384))*SQRT(SUMSQ(B1:B384)))`

`=SQRT(SUMSQ(A1:A384))`

`=ROWS(A1:A384)`

0.5111: top_k's score, within 1e-12

0.5111, within 1e-12

1.0000: length 1 within 1e-6

384

No tokenizer or sentence model in Excel: the 384 numbers come from MiniLM.


Save the day: show the commit

```
git show --stat HEAD
test_textkit.py | 69 ++++++
textkit.py      | 134 ++++++
2 files changed, 203 insertions(+)
```

pytest: 37 passed. textkit only grows: nothing removed.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: the loop's range drops the last partial chunk: 347 chunks, 53 sections with none, 4.12's clause (e) lost.

Code that runs is not the same as code that is right.

What you can do now

- Chunk sections, never across a boundary
- Load a pinned model; know its cache
- Embed once; dot product equals cosine
- Search by meaning; report section, score, words
- Count word pieces against the 256 limit

Week 2, Day 2 Zero-shot classification against TF-IDF