



Python for Fixed Income

COURSE 7 · NLP AND LLMS

Reading EDGAR

WEEK 2 · DAY 4

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

Today: reading EDGAR

Three documents, as filed

10-K	fiscal 2025, 4,008,734 bytes
Exhibit 4.1	the Boyd indenture, 2021
8-K	two sections of one filing

Today

1. EDGAR's addresses, as strings
2. The SEC's rules for programs
3. `sec_headers` and `edgar_get`
4. HTML to text; inline XBRL
5. A 10-K cut into its items
6. The indenture and an 8-K
7. Excel, Git, the AI check

`FETCH_LIVE = False`: the notebook sends no request; it reads the course's copies.

The rules for real documents

Today's numbers: sizes, positions, and word counts of documents

Allowed names

a word count of an item

what the document states

the filer's own heading

Never called

a measure of risk

a view on the company

a comparison of issuers

Every excerpt has its source line. Company names appear only there, in the data lines, and in the addresses, as authors.

EDGAR's addresses

906553

CIK: the SEC's number for the filer

0001437749-26-004908

accession number: one filing

```
https://data.sec.gov/submissions/CIK0000906553.json  
https://www.sec.gov/Archives/edgar/data/906553/000143774926004908/  
https://www.sec.gov/Archives/edgar/data/906553/000143774926004908/index.json  
https://www.sec.gov/Archives/edgar/data/906553/000143774926004908/bgc20251121_10k.htm
```

The CIK padded to ten digits in one, plain in another; dashes dropped in the folder. Build, never retype.

"EDGAR started in 1994/1995."

One line, word for word, on the SEC's page "Accessing EDGAR Data", the page that asks programs to declare a user agent. The page says no more about it.

2021	the Boyd indenture, Exhibit 4.1 to an 8-K
-------------	---

2019 to 2025	the 8-K sections in the course's file
---------------------	---------------------------------------

2026	the Boyd 10-K for fiscal 2025
-------------	-------------------------------

Source, read 8 October 2026: sec.gov/search-filings/edgar-search-assistance/accessing-edgar-data

The SEC's rules for programs

A declared user agent

"Please declare your user agent in request headers": a name and an email.

At most 10 requests a second

"... no more than 10 requests per second, regardless of the number of machines"

Copies, cited

Public information, "may be copied or further distributed"; cite the SEC as the source.

Sources, read 8 October 2026: sec.gov/search-filings/edgar-search-assistance/accessing-edgar-data; sec.gov/about/privacy-information

`FETCH_LIVE = False` in the notebook. Your own name and email go in your own script.

sec_headers and edgar_get

```
declared = textkit.sec_headers("Student Name", "student@example.com")  
body = textkit.edgar_get(url, declared)    # sec.gov hosts only; 0.11 s apart
```

sec_headers

No default contact. One @ with a . after it, or ValueError.

Retried

429 and 503: wait 1, 2, 4, 8 seconds, then try again.

Never retried

403, 404, and the rest: RuntimeError with the status.

The tests use a fake session and a fake clock: no network, no waiting.

Inline XBRL, and what html_to_text keeps

4,008,734

bytes of HTML, as filed

387,293

bytes of the hidden ix:header

2,296

tagged facts in the file

422,118

characters of text

```
soup.find_all(["script", "style", "ix:header"]) # removed, with all they hold
```

The text is 0.1053 of the file's size, 65,067 words. The header holds passages the visible text holds too: kept, they would count twice. A served copy, with sec.gov's script tag, gives the same text: True.

A block tag starts and ends a line; a cell is one piece of text; inline tags add nothing.

A 10-K cut into its items

Item	Words
1	7,312
1A	5,446
1C	1,175
7	14,826
7A	573
8	29,178
15	2,041

Two traps

The table of contents: 46 lines start with an item code, 23 of them before Item 1. "ITEM 1A. Risk Factors 11" at character 4,605; the heading at 53,289.

Mentions inside a paragraph: 9 of 55, such as "See Item 8 of Part II". Never cut.

23 of 23 items found: a heading starts a line, and items are found from the last to the first.

The indenture: the same text, character for character

872,436

bytes: boyd_2021.htm, as filed

392,021

characters of html_to_text

63,595

words

True

equal to boyd_2021.txt

```
assert textkit.html_to_text(exhibit_bytes) == Path("course7_indentures/boyd_2021.txt").read_text(...)
```

Every number about the indenture since day 1 came from this function on this exhibit.

Source: Boyd Gaming Corporation, indenture for the 4.750% Senior Notes due 2031, Exhibit 4.1 to the 8-K filed 2021-06-08, accession 0001193125-21-185495.

8-K items: a cross-reference must not cut

```
Item 1.01 Entry into a Material Definitive Agreement
...the stored body...
Item 2.03 Creation of a Direct Financial Obligation ...
The information contained in Item 1.01 ... in this Item 2.03.
```

157 and 21

words: each body equal to the stored one

5

pieces from a split at every mention

Only a heading that starts a line cuts. The filer's cross-reference stays inside its sentence.

Excel side by side

No HTML parser, no inline XBRL, no request to sec.gov in Excel. Cutting pasted text, yes.

Excel side by side (tested in Excel)

```
=TEXTBEFORE(TEXTAFTER(A1,"ITEM 2. Properties"&CHAR(10)),CHAR(10)&"ITEM 3.")  
=ROWS(TEXTSPLIT(TRIM(SUBSTITUTE(B2,CHAR(10)," ")),," "))  
=TEXTAFTER(A1,CHAR(10)&"ITEM ",-1)  
=TEXTBEFORE(TEXTAFTER(A1,"ITEM 3. Legal Proceedings"&CHAR(10)),"Item ")  
=TEXTAFTER(A1,"ITEM 7.")
```

Item 2's body, equal
242
after the last heading
"See"
#N/A

One cell, 1,704 characters: the whole text, 422,118, is over a cell's 32,767.

Save the day: a change, word by word

```
test_textkit.py | 97 ++++++
textkit.py       | 147 ++++++
```


git diff --word-diff textkit.py, after 0.11 became 0.05:

```
... min_interval: float = [-0.11,-]{+0.05,+} retries: int = 4) -> bytes:
```

git restore textkit.py puts the file back. pytest: 49 passed.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: the helper sends User-Agent "python-requests" and never uses the caller's name and email.

Code that runs is not the same as code that is right.

What you can do now

- Build EDGAR's addresses from a CIK and an accession
- Keep the SEC's rules in code: a contact, 10 a second
- Turn HTML into text, without the XBRL header
- Cut a 10-K at its items, past the contents
- Cut one item from pasted text in Excel

Week 2, Day 5 Case study: where the models disagree