



Python for Fixed Income

COURSE 7 · NLP AND LLMS

First LLM Request

WEEK 3 · DAY 1

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

Today: first LLM requests

Two documents

FOMC	the statement of 2026-09-16
Indenture	Boyd, Section 3.07(b), one paragraph
Saved	responses recorded 2026-10-08

Today

1. A request, and OpenRouter
2. The key, kept out of Git
3. ask: one door for every request
4. First request, checked by regex
5. Tokens, cost, reasoning off
6. Temperature; replay; failures
7. OpenAI, Excel, Git, AI check

Every response in the notebook is a saved response, replayed. A live run can differ.

The rules for real documents

Every prompt asks what a document says: a quote, a summary, or NOT IN TEXT

Allowed names

the response

the saved response

the model returned

Never called

the model's opinion or view

what the model believes

what the model knows

The system prompt goes first in every request and refuses an opinion, a grade, a comparison of issuers, or a view about the future.

What a request is

```
{  
  "model": "nvidia/nemotron-3-super-120b-a12b:free",  
  "messages": [  
    {"role": "system", "content": SYSTEM_PROMPT},  
    {"role": "user", "content": question + text}],  
  "temperature": 0, "seed": 42, "max_tokens": 2000  
}
```

What comes back

content	the response's text
model	the id that answered
usage	tokens in and out, billed
finish_reason	stop, or length

The model sees only the messages: not your file, not your folder, not your machine.

OpenRouter: one key, many models

20

free-model requests a minute

50

a day, under 10 credits bought

1,000

a day, 10 credits or more

Key	Model	Price per million tokens
text	nvidia/nemotron-3-super-120b-a12b:free	US\$0.00 in, US\$0.00 out
fallback	google/gemini-2.5-flash-lite	US\$0.10 in, US\$0.40 out
vision	google/gemini-2.5-flash-lite (day 13)	the same

Sources, read 8 October 2026: openrouter.ai/docs/api-reference/limits; openrouter.ai/api/v1/models

The key

get_secret looks in order

1. the environment
2. Colab's Secrets panel
3. my_bondmath\.env

Never

print the client (client.api_key shows it)

paste it in a cell, a prompt, a command

commit .env

```
!git -C "{work}" check-ignore -v .env cache/llm/  
.gitignore:4:.env      .env  
.gitignore:5:cache/    cache/llm/
```

The key cell prints only: OpenRouter key loaded: False. Mode: replay.

ask: one door for every request

1 System prompt first

quote exactly, give the section, NOT IN TEXT, refuse opinions

2 A fingerprint

request_key: SHA-256 of provider, model, messages, parameters

3 The cache

one JSON file per key; replay returns it or raises LookupError

4 Refusals

429 and 5xx waited and retried; others stop with a plain message

No header and no key in the fingerprint or the file. pytest at the end of the day: 59 passed.

The first request

The Committee decided to raise the target range for the federal funds rate by 1/4 percentage point to 3-3/4 to 4 percent, in support of the Federal Reserve's dual mandate.

```
Model nvidia/nemotron-3-super-120b-a12b:free, temperature 0, seed 42; saved response recorded 2026-10-08; replayed (no key loaded).
```

3-3/4 to 4 percent

in the statement, by day 1's regex

3-3/4 to 4 percent

in the saved response

yes

the same range

A number a model returns is checked against the document by code that does not involve the model.

The system prompt on one paragraph

Question	Saved response
Share redeemable with Equity Offering proceeds, and the price?	40%, 104.750%: each in the paragraph
What is the interest rate on the Notes?	NOT IN TEXT

Section 3.07(b), 513 characters. The rate is in the indenture's title, not in the paragraph sent.

Source: Boyd Gaming Corporation, indenture for the 4.750% Senior Notes due 2031, Exhibit 4.1 to the 8-K filed 2021-06-08, accession 0001193125-21-185495, Section 3.07.

Tokens and cost

323 / 43

prompt / completion tokens, free model

US\$0.00000000

on the free model

US\$0.0000495

the same counts at the fallback's price

308 / 38

the fallback's own count of the same text

US\$0.0000460

usage_cost

US\$0.0000460

reported by OpenRouter

$$\text{cost} = \text{prompt_tokens} \times \text{prompt price} + \text{completion_tokens} \times \text{completion price} \quad \# \text{ US dollars}$$

Reasoning switched off

43

completion tokens, reasoning off

1261

completion tokens, reasoning on

1222

of them reasoning, never shown

```
extra_body=textkit.NO_REASONING    # {"reasoning": {"enabled": false}}
```

Same sentence quoted either way. The answer is in the text, so the course asks the model to find and quote; reasoning tokens are completion tokens, and a paid model bills them.

Temperature and seed

```
t0 s42 Inflation remains elevated, and today's policy action will support a timelier return to ...  
t0 s43 Inflation remains elevated, and today's policy action will support a timelier return to ...  
t0 s44 Inflation remains elevated, and the Committee's policy action is intended to support a t...  
t1 s42 Inflation remains elevated, and the rate increase aims to support a timelier return to t...  
t1 s43 Inflation remains elevated, but the Committee's policy action will support a timelier re...  
t1 s44 Inflation remains elevated, but today's policy action will support a timelier return to ...
```

2

distinct texts of 3 at temperature 0

3

distinct texts of 3 at temperature 1

Temperature 0 and a seed narrow the spread. They do not promise identical text.

Replay, and when a request fails

```
cache/llm/0e28d64e7792....json (14 saved responses: request, content, usage, recorded; no header, no key)
two words added, replay: LookupError: no saved response for request dd2ca18e0b82 in llm
```

```
a wrong model id (HTTP 400): The provider refused the request with HTTP 400
a refused key (HTTP 401): The provider refused the key: check OPENROUTER_API_KEY in your .env
no credit on a paid model (HTTP 402): The provider refused the request with HTTP 402
```

- 400** bad request: a wrong model id
- 401** the key
- 402** no credit
- 429** too many: a wait can clear a per-minute limit, not a daily cap

The OpenAI route

```
PROVIDERS = {"openrouter": {"base_url": "https://openrouter.ai/api/v1", "key_name": "OPENROUTER_API_KEY"},  
            "openai": {"base_url": "https://api.openai.com/v1", "key_name": "OPENAI_API_KEY"}}
```

gpt-4.1-mini-2025-04-14

the model id OpenAI returned

None

provider: an OpenRouter field

The provider is part of the fingerprint: one saved response per route.

Excel and Git

Excel side by side (tested in Excel)


<code>=B2*C2+B3*C3</code>	0.000046
<code>=SUMPRODUCT(B2:B3,C2:C3)</code>	0.000046
<code>=TEXT(D1,"0.0000000")</code>	0.0000460
<code>=B2*E2+B3*E3 (prices as text)</code>	0.000046
<code>=SUMPRODUCT(B2:B3,E2:E3) (prices as text)</code>	0

```
!! __pycache__/  
!! cache/
```

SUMPRODUCT treats text cells as zero. git status --ignored: cache/ marked !!, out of Git.

The AI check

The docstring is the prompt. The code that comes back runs, and its number looks fine.



Today: the helper prints every header, and Authorization holds the whole key.

Code that runs is not the same as code that is right.

What you can do now

- Say what a request holds and what comes back
- Keep the key in .env, out of Git, never printed
- Send through ask; check numbers with code
- Price tokens with usage_cost, free or paid
- Replay saved responses; read 400 to 429

Week 3, Day 2 Covenants 1: the anatomy of an indenture