pfi

Python for Fixed Income

COURSE 8 · IMPLEMENTING AGENTIC AI

Structured Answers and Traced Numbers

WEEK 1 · DAY 4

Created by **Jeff Lenamon**

This course is educational content created in a personal capacity. Nothing here is investment advice or a recommendation to buy or sell any security. All portfolio data is synthetic.

Today: structured answers and traced numbers

What day 3's saved answers did

2 of 6 wrote the tool's number as returned

4 of 6 rounded it, unasked

147 bp copied from a tool description

Today

1. Free text against a schema
2. Figure and DeskAnswer
3. A structured run
4. check_figures
5. A wrong digit, no tool behind it
6. numbers_untraced
7. Traces, and still wrong

Every model response in the notebook is a saved response, replayed. A live run can differ.

The rules for agents

What an agent produces: text, every number from a tool, now checked in code

Allowed names

the agent returned

the saved run

the brief states

Never called

the agent's opinion or view

what the agent believes

the agent's advice

On the synthetic book: a description, never a recommendation

Allowed names

DV01 per 100 of face

position DV01 in dollars

the largest DTS contributions in the file

Never called

cheap, rich, or attractive

a trade idea or a recommendation

overweight or underweight

Free text against a schema

The saved free-text answer to day 3's duration question

```
"As of 2026-09-30, the modified duration ... of bond 99126WAE0, ...

... at a 4.516% yield, is approximately 2.588 years."
```

Model gpt-4.1-mini-2025-04-14, temperature 0, seed 42; 2 requests, 2,238 prompt and 125 completion tokens; saved run recorded 2026-10-11; replayed (no key loaded).

find it

the number, in a sentence

match it

to a tool result, by eye

count

the decimals it claims

Nothing in the sentence says which tool request a number came from.

Figure and DeskAnswer

```
type: json_schema; name: DeskAnswer; strict: True
DeskAnswer: required ['text', 'figures', 'sources'], extra fields refused
 text string the answer in plain sentences
 figures array every number the text states
 sources array the files, sections, or statement dates used
Figure: required ['name', 'value', 'unit', 'tool_call_id'], extra fields refused
 name string what the number is, for example 'DV01 of the 2031 bond'
 value number the number, written to the decimals stated in the text
 unit string the unit as the tool result's key ends: per_100, pct, bp, usd, years
 tool_call_id string the id of the tool request whose result holds the number
```

answer_format: a strict JSON schema. Every field required, no extra field, descriptions sent as prompt.

A structured run

The same question with response_format=STRUCTURED: in the saved run

figure	`Modified duration of bond 99126WAE0: 2.587515 years`
cites	`call_ReR4zEhA1ezTVw3Wji8NBh5a`
run's ids	`'call_ReR4zEhA1ezTVw3Wji8NBh5a'`
text says	`"... approximately 2.588 years ..."`

Model gpt-4.1-mini-2025-04-14, temperature 0, seed 42; 2 requests, 2,802 prompt and 194 completion tokens; saved run recorded 2026-10-11; replayed (no key loaded).

The figure carries the tool's number whole; the sentence still rounds it. Check both.

check_figures: cited, at its rounding, in its unit

Three questions per figure

1. Is the cited request in this run?
2. Does its result, rounded to the figure's own decimals, equal the figure?
3. Does that number's key end in the unit?

2.588 for 2.587515 passes; a wrong digit fails

In the saved runs

```
duration  1 of 1 passed
price     1 of 1 passed
dts       2 of 2 passed
```

No judgement and no reading of the text: an id, a number, a unit.

A wrong digit, and no tool behind it

Two copies of the duration answer, checked against the same run

copy	value	cited	value_matches	passed
one digit changed	2.587516	True	False	False
no tool behind it	2.587515	False	False	False

What the agent did, step by step

```
1  assistant dts  {"spread_duration_years":0,"spread_bp":0} tool_calls 1343 22
1  tool dts  {"dts_years_bp": 0.0, "tool_call_id": "call_yIfwkj8f3hpnKWdtIRUwORkp"} <NA> <NA>
```

The OAS question, structured: no figure in the answer (figures: 0); "cannot be determined"

A figure with no tool behind it has nothing to be checked against, so it cannot pass.

numbers_untraced: checking free text

Every number in the text must trace, at its rounding, to some tool result

Skipped by rule: ISO dates

Skipped by rule: the unit phrase "per 100"

Skipped by rule: every number the question wrote

```
untraced in the saved answer: []
untraced after one digit changed: ['2.589']
```

No citations, so no lookup: a number equal to some other result passes. check_figures is stricter.

A figure that traces can still be the wrong figure

Two bonds' DV01s; values and citations swapped, names kept

check_figures after the swap: 2 of 2 passed

```
DV01 of bond 99039NAC0: 0.127083 against the key 0.032627: differs

DV01 of bond 99033HAF2: 0.032627 against the key 0.127083: differs
```

Model gpt-4.1-mini-2025-04-14, temperature 0, seed 42; 3 requests, 4,487 prompt and 327 completion tokens; saved run recorded 2026-10-11; replayed (no key loaded).

The right tool on the wrong bond passes every check here. Day 13 adds answer keys decided first.

Excel and Git

Excel side by side (tested in Excel)

`=ROUND(B2,D2)=C2`	`B 2.587515  C 2.587515  D 6` (duration, figure)	TRUE
`=ROUND(B3,D3)=C3`	`B 2.587515  C 2.588  D 3` (duration, text)	TRUE
`=ROUND(B13,D13)=C13`	`B 2.587515  C 2.587516  D 6` (planted wrong digit, figure)	FALSE
`=ROUND(B15,2)=ROUND(C15,2)`	`0.0437 against 0.043625` (the AI check's bug)	TRUE
`=ROUND(B15,4)=C15`	`at the figure's own 4 decimals`	FALSE

Excel against the course's checks: 12 of 12 rows the same

```
-free,gpt-4.1-mini-2025-04-14,2,2238,125,0.001095,2,answer,,
+free,gpt-4.1-mini-2025-04-14,2,2238,125,0.001095,2,answer,0,True
```

runs.csv: 6 rows, figures_traced and passed filled; git diff shows only those. 24 passed.

The AI check

The docstring is the prompt. The agent code that comes back runs and looks right.

1

Docstring

Round the tool value to the figure's own decimals. 0.0437 fails.

→

2

Code back

Written for the course in the style of an AI answer. Read it.

→

3

Test first

0.0436 must match 0.043625 and 0.0437 must not. Run it: it fails.

→

4

Fix

Read the decimals from the figure; round only the tool value.

Today: both sides rounded to 2 decimals: 0.0437 and 0.043625 both become 0.04, so a wrong digit passes.

Code that runs is not the same as code that is right.

What you can do now

- Ask for an answer as a strict JSON schema
- Every figure names its tool request
- check_figures: cited, rounded, in its unit
- numbers_untraced for free text, exceptions stated
- A traced figure can still be the wrong bond

Week 1, Day 5 Case study: a risk agent on the synthetic book